

# Convolutional neuronal networks

Till Korten

With Material from

Johannes Müller, Robert Haase: PoL



Alex Krull, Uwe Schmidt: MPI CBG


Martin Weigert: EPFL Lausanne

Ignacio Arganda-Carreras: Universidad del Pais Vasco


December 2022

- Neural networks are composed of individual artificial “neurons”

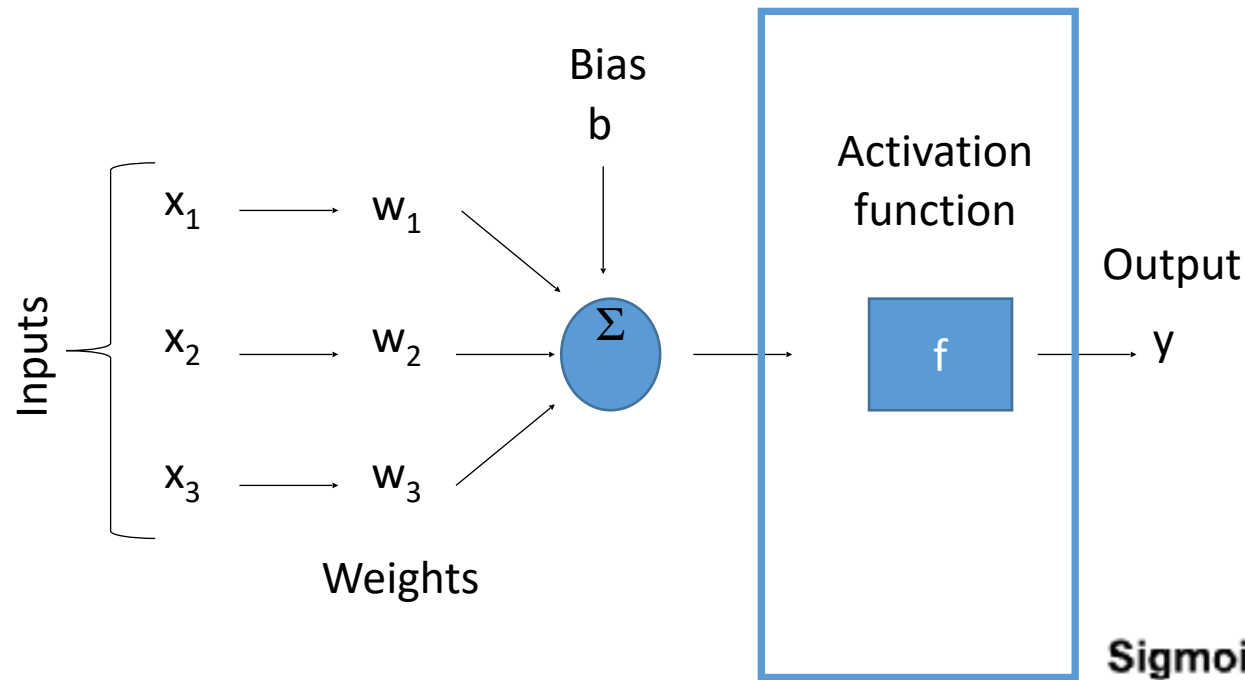




- Neural network: multiple layers of artificial “neurons”
- Deep neural network (DNN): neural network with more than one layer between input and output



# An artificial “neuron” is a mathematical function



Single neuron output calculation  
 $y = w_1x_1 + w_2x_2 + w_3x_3 + b = w^T x + b$


For image data, the values  $x_1, x_2, \dots$  would be

Pixel intensities


Pixel coordinates

**Sigmoid**

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




**tanh**  
 $\tanh(x)$



**ReLU**  
 $\max(0, x)$




**Leaky ReLU**  
 $\max(0.1x, x)$




**Maxout**  
 $\max(w_1^T x + b_1, w_2^T x + b_2)$


**ELU**  
$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




# Network layers can have different architectures




Fully connected layer



Convolutional layer



Pooling layer




Convolutional layer

**Previously:**

Defined filter kernels

$1/16$	$1/8$	$1/16$
$1/8$	$1/4$	$1/8$
$1/16$	$1/8$	$1/16$




**Now:**


Learned filter kernels

$w_{11}$	$w_{12}$	$w_{13}$
$w_{21}$	$w_{22}$	$w_{23}$
$w_{31}$	$w_{32}$	$w_{33}$


**Input**



**Output**




# Pooling layers reduce the layer size



Pooling layer

3	15	1	13
9	7	0	10
11	5	5	3
1	8	9	6



Maximal values

15	13
11	9

Average values


8.5	6.0
6.3	5.8

# The network learns by minimizing a loss function

- Learning is an optimization problem
- Step 0: Initialize the network randomly
  - Weights
  - Bias
- Step 1: Forward pass the input through the network, get an initial prediction (Images 0...M)
- Step 2: Compare the output with the ground truth, computer the error (loss function)
  - The **loss function** can be freely defined.
  - Mean squared error:

$$\mathcal{L}(y, \hat{y}) = \frac{1}{M} \sum_{i=1}^M (\hat{y}_i - y_i)^2$$

- Step 3: Update weights





The loss function can be expanded from

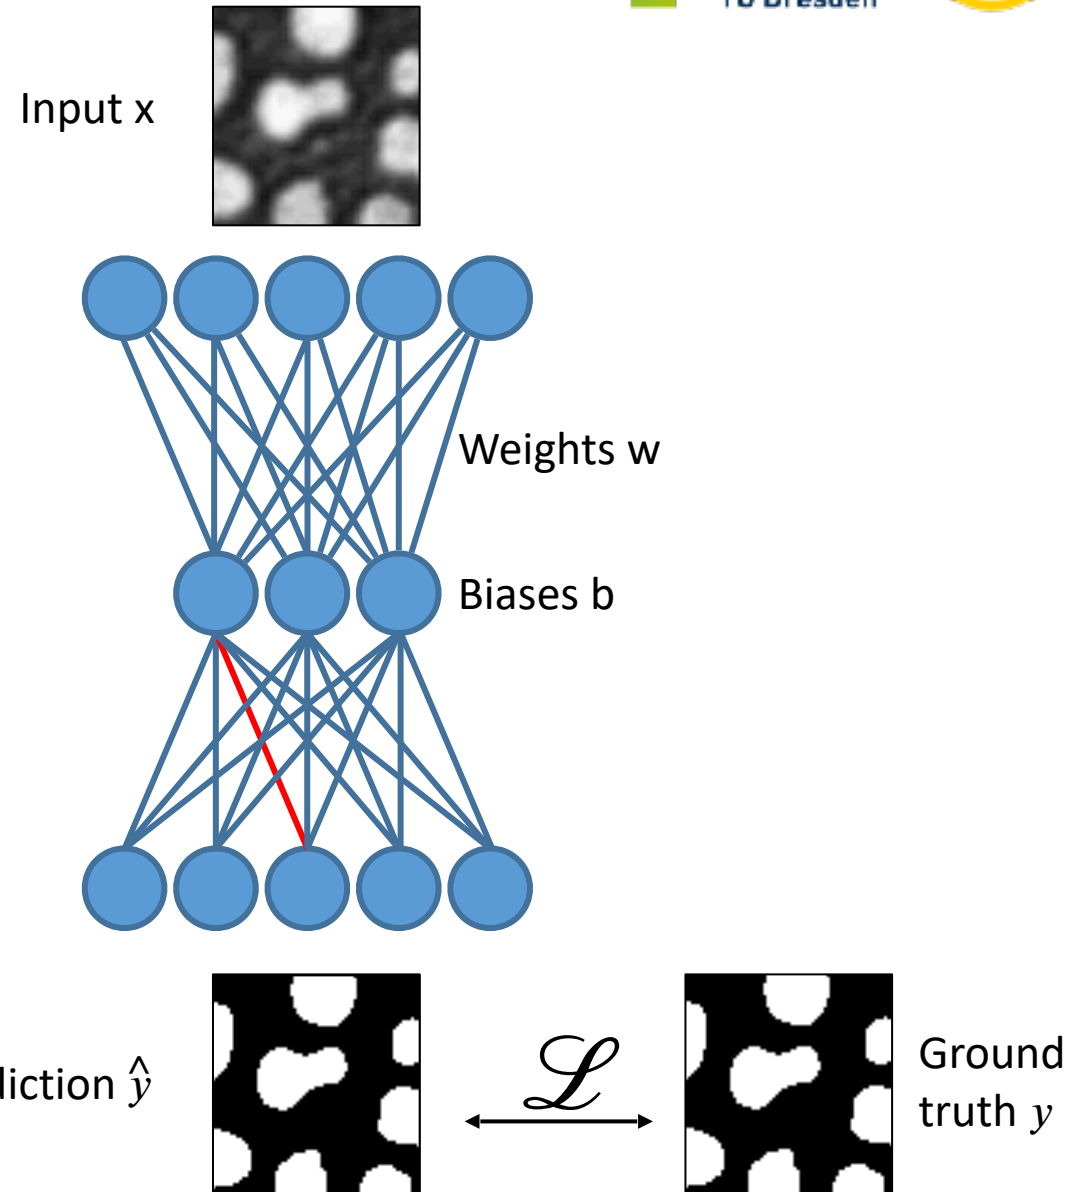
$$\mathcal{L}(y, \hat{y}) = \frac{1}{M} \sum_{i=1}^M (\hat{y}_i - y_i)^2$$

as the prediction depends on inputs  $x$  weights  $w$  and bias  $b$

We can calculate derivatives with respect to  $w$  and  $b$  to find their optimal values

→ Derivatives tell us how to change  $w$  &  $b$  in order to improve the prediction (i.e. minimize the loss function)

Repeat this for each layer, update weights  $w$



## Popular frameworks:

<https://www.tensorflow.org/>

<https://www.pytorch.org/>

**Hardware requirements:** Nvidia (CUDA-capable) graphics card (GPU)

**Memory:** The more GPU memory the better



[Google LLC](#), Public domain, via Wikimedia Commons



[PyTorch](#), [BSD](#), via Wikimedia Commons



[Adam Kapetanakis](#), [CC BY-SA 4.0](#), via Wikimedia Commons

## Before you start training:


- Split data into three groups:
  - Training set
  - Testing set
  - Validation set

## During training:

- Apply network to training set:
  - Measure performance (“loss”) of network, update weights
- Apply network to testing set:
  - Measure performance of updated network, keep weights unchanged
- Improve network architecture (optimize “hyperparameters”)





## During validation:

- Measure performance in unseen validation dataset



### Overfitting:



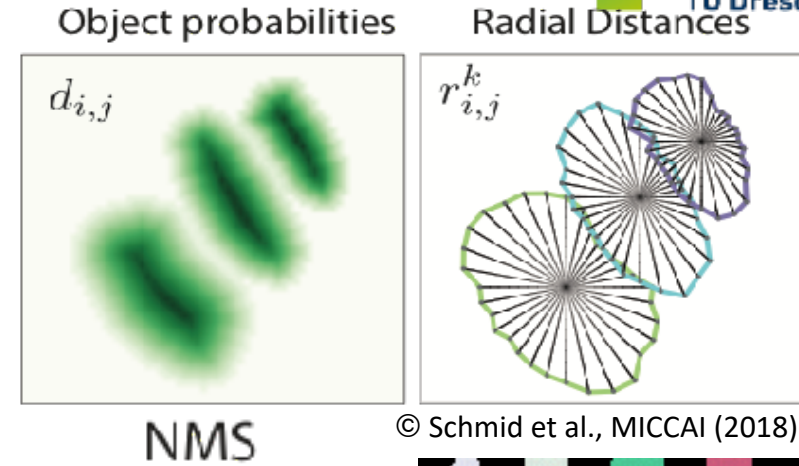
- Network is learning things „by heart“
- Hint at this happening:  
Updated weights from training fail to perform well in test




- The **U-net** is the most used network architecture in biological image processing using CNNs.
  - Encoder: Increase the “What”, decrease the “Where”
  - Decoder: Use the “What”, to identify the “Where”

## Strategy:

- Add additional information to prediction
- Member pixels of objects (nuclei) can be reached via a straight line from the center





Dense Polygon Prediction  
(e.g. U-Net, ResNet)




Polygon Selection  
(Non-Maximum Suppression NMS)



Noisy images + Crowded cells = Common source of segmentation errors





Dense Segmentation  
(e.g. U-Net)



Bounding box based methods  
(e.g. Mask-RCNN)

Object probabilities

Radial Distances











© Schmid et al., MICCAI (2018)



## Non-maximum-suppression (NMS):

- Object probabilities: Probability that pixel belongs to class “nucleus”
- Multiple maxima lead to multiple possible polygons for the same nucleus

## Algorithm:

- Select polygon with highest object probability inside: 
- Look at other polygons: Is the overlap of  with  larger than threshold  $\tau$ ?
  - Yes:  and  are actually the same object, drop 
  - No:  and  are separate nuclei
- Setting  $\tau$  very high leads to many false positives!


## Non-maximum suppression




© Schmid et al., MICCAI (2018)




- **CARE**: content-aware restoration
- Image acquisition of pairs of images: A high-quality and a low-quality image.
- Caveats:
  - Reconstructs shot noise present in high quality training images
  - Trained model only applicable to image data of the same conditions (biological sample, microscope, etc)



5 example validation patches  
top row: input (source), middle row: target (ground truth), bottom row: predicted from source



© MPI-CBG, Dresden. All rights reserved.



Raw data = signal + noise



## Strategy:

- Try to predict intensity of pixel  $y$  from surrounding pixels  $x$
- CNN fails to predict noise component → N2V can only reproduce signal from the surroundings of  $y$

## Beware:

- Only **random** noise can be removed, otherwise artifacts occur


<https://github.com/juglab/n2v>

<https://forum.image.sc/t/n2v-artefacts-in-training-data/70686>

- Images are tiled
- limited “receptive field of the network”

## **Receptive field:**

- Objects must be smaller than receptive field to be detectable



## Unbalanced training data:

- Some labels appear more often in training data than others
- Rare events will not be learned because missing them doesn't harm accuracy much
- Weighted data sampling
- **Biased results!!**



### Is the iPhone racist? Chinese users claim iPhoneX face recognition can't tell them apart

APPLE has come under fire following numerous complaints from Chinese users who claim the iPhone X face recognition can't tell them apart.


<https://www.news.com.au/technology/gadgets/mobile-phones/is-the-iphone-racist-chinese-users-claim-iphonex-face-recognition-cant-tell-them-apart/news-story/13814540e8c82ad466aca687e12af64c>



# When the input data does not fit to the training data



Prediction






What happened here?

Receptive field too small

I used a different resolution than during training

Overfitting



- With great power comes great responsibility: **Validate your models well!**
- Better data more important than better model
- Often performs fantastic – *but you don't know why*
- Generative neural networks (like CARE) can dream up data – to a hammer everything looks like a nail!