

# How to Teach “Convolutional Neural Networks”



**Stephanie Käs**  
*M. Sc. Physics*

**HybridLaunch Gießen n**

*Hybrid Rocket Developme*

**Speaker & Marketing Manager & Graphic**

**AI Center RWTH Aac**

*Computer Vision & Machine*

**Research & Teaching A**

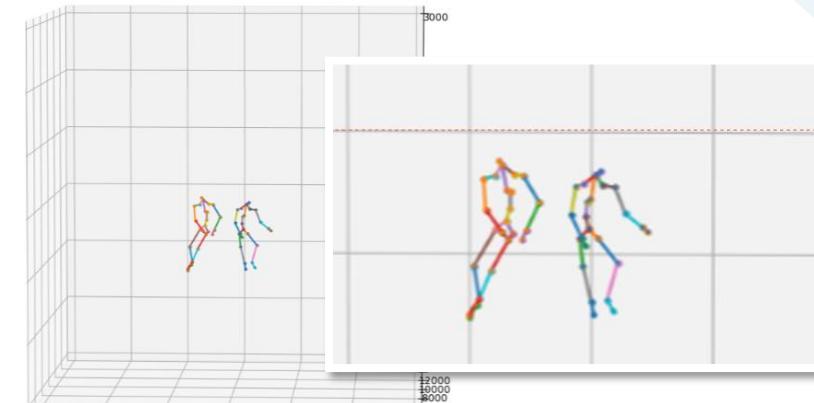
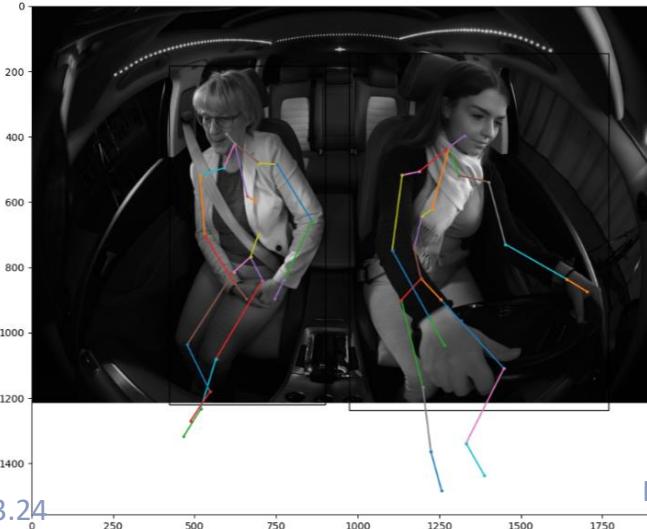
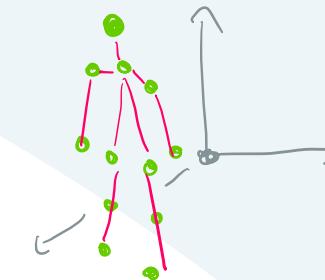
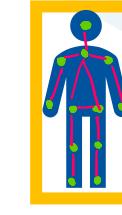
**JLU Gie**

*Statistik für Ge*

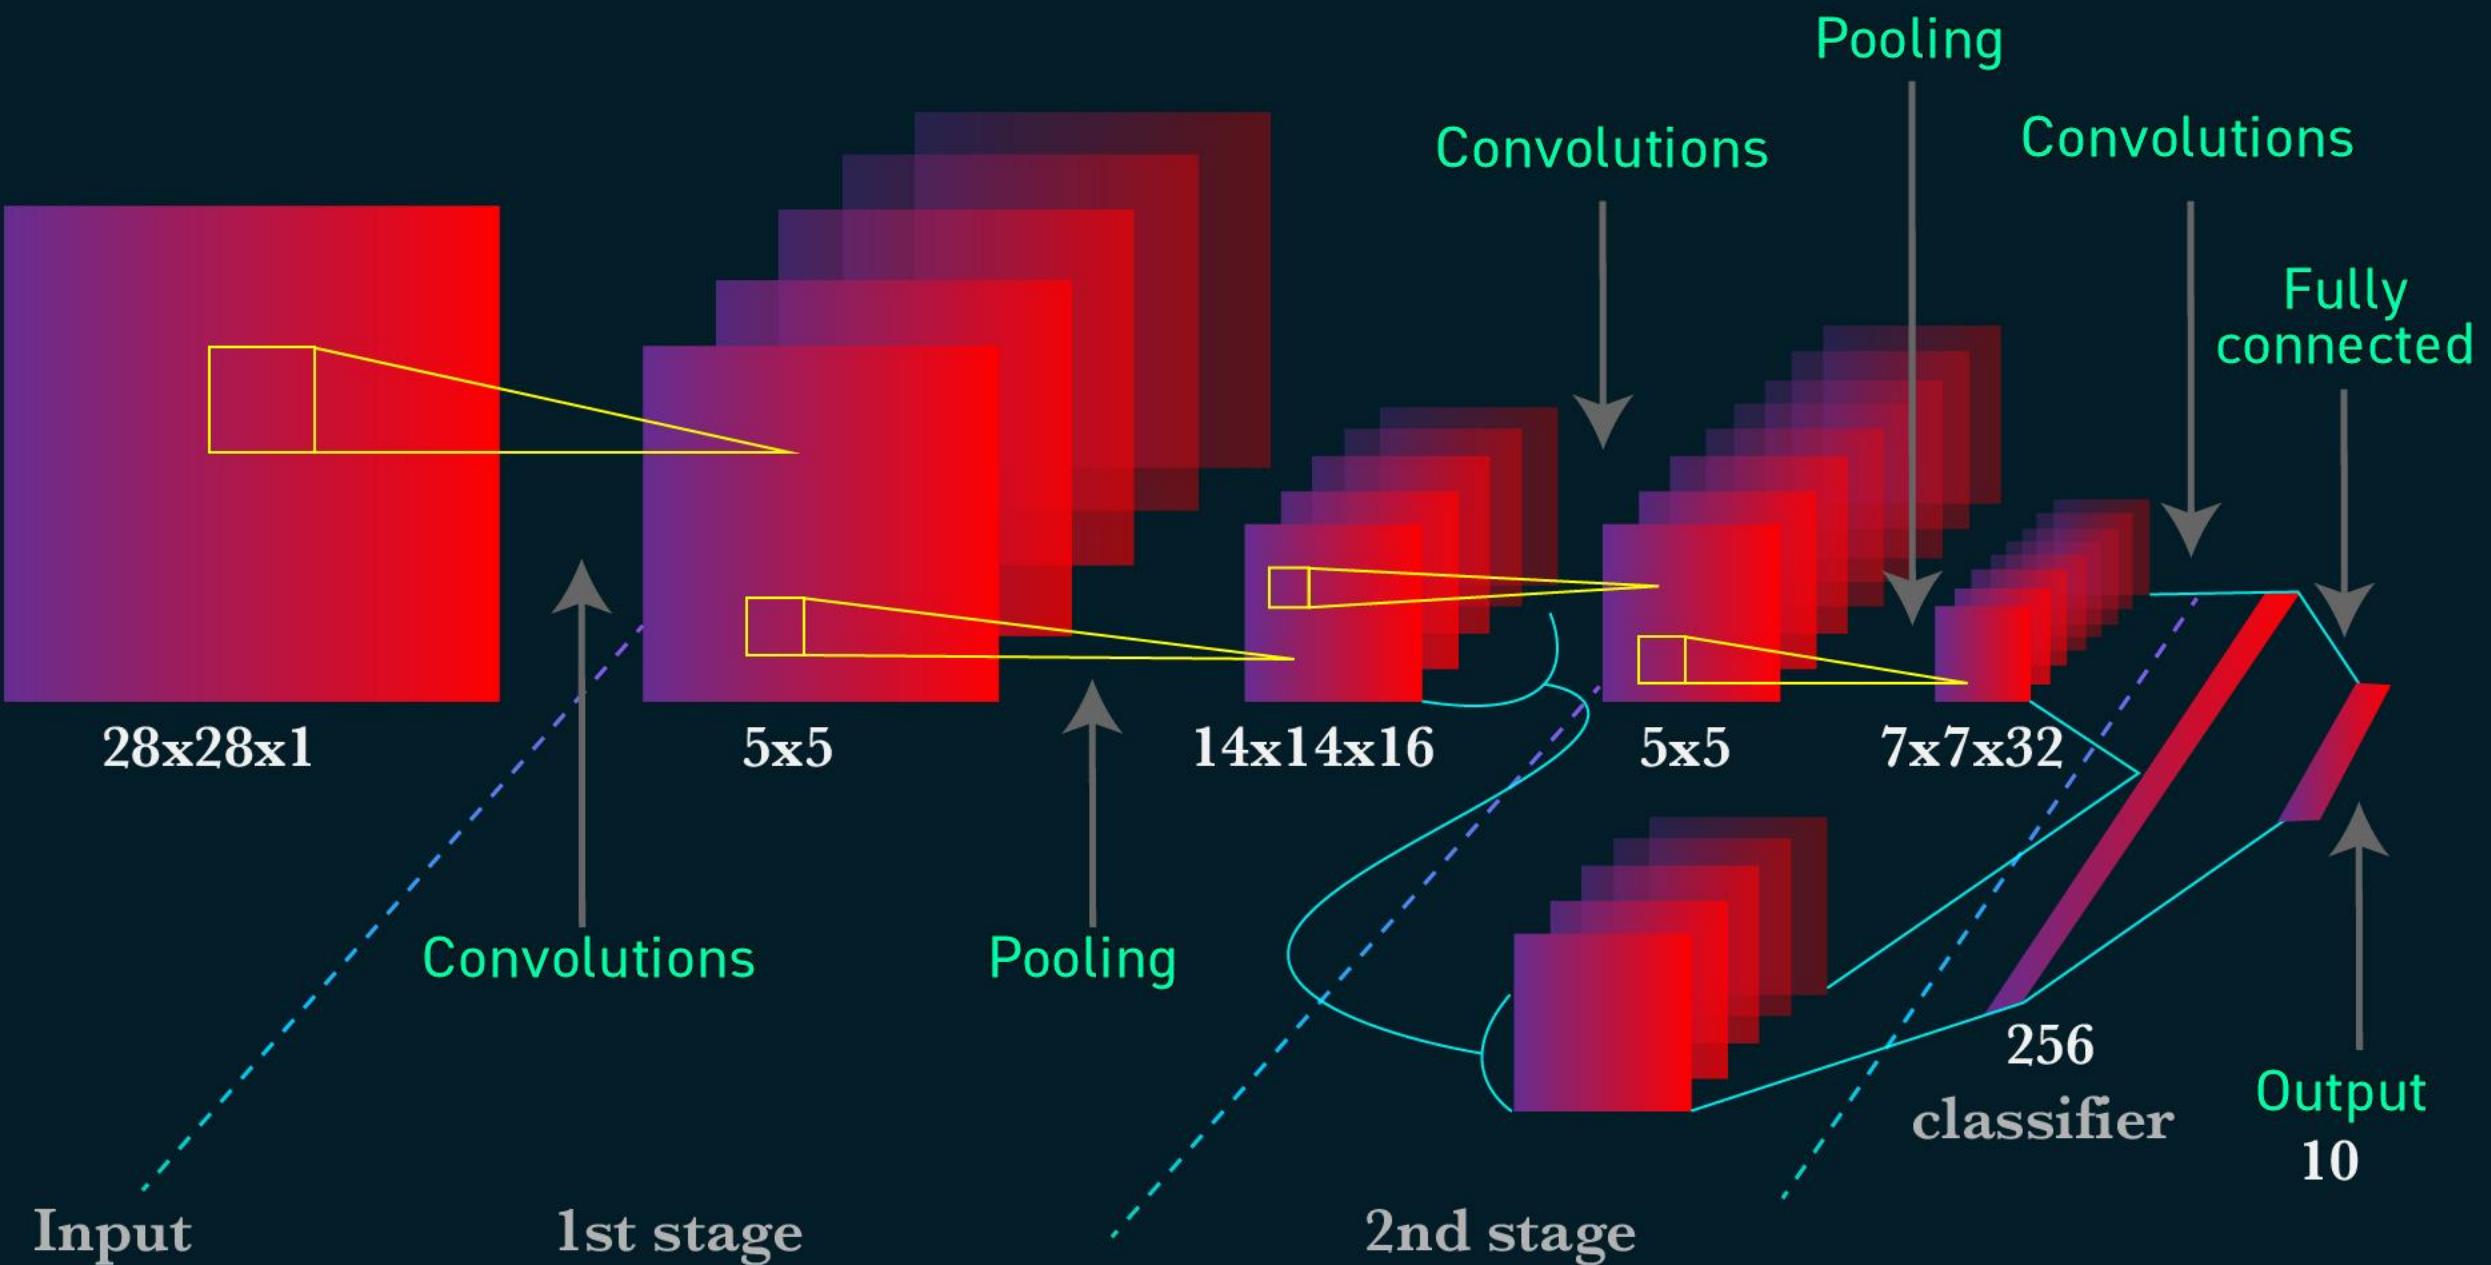


# Current Research

- Computer Vision Group @ RWTH Aachen University
- Estimating Poses of Humans from Images
- Processing images using Convolutional Neural Networks



# Convolutional Neural Networks - Recap



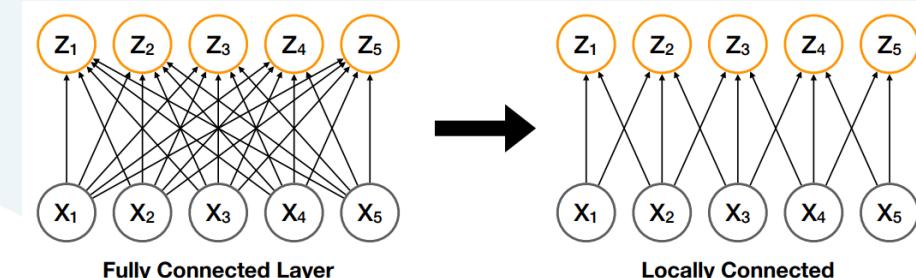
# What might be hard to understand?



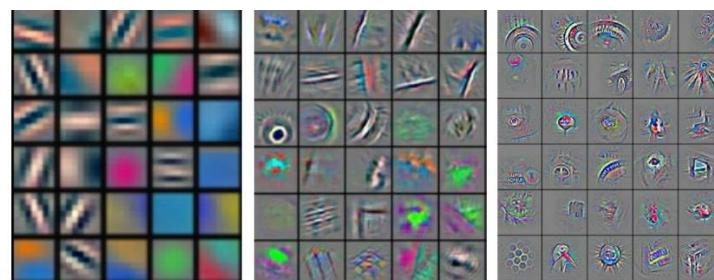
Images as input



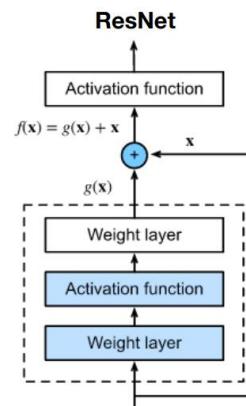
Concept of channels



Reasons for parameter reduction



What filters actually learn



Issues during the training process

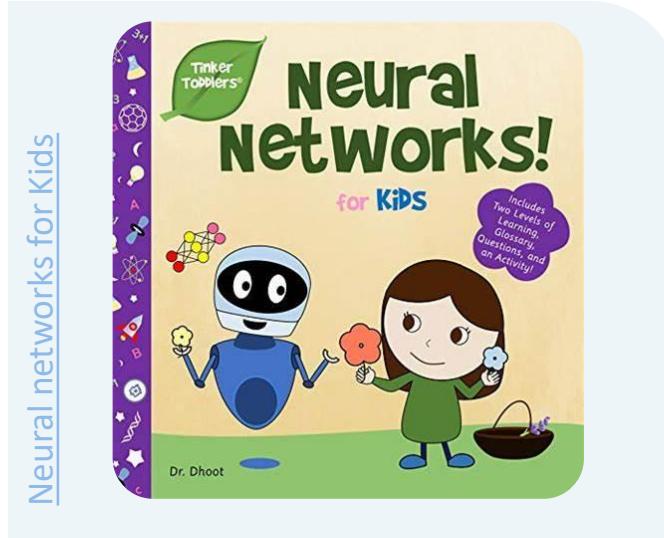
Regularization & training tricks

How to set hyperparameters?

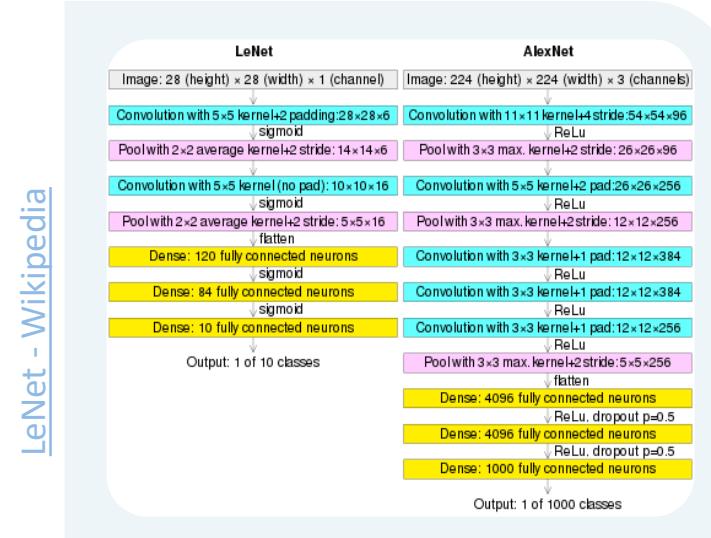
# What might be hard to understand?



# How can we teach CNN basics?



Explain general concepts



LeNet - Wikipedia

Explain common architectures

Databricks Blog

```

1 import torch
2 import torch.nn as nn
3
4
5 class TwoLayerNet(nn.Module):
6     """
7         In the constructor we instantiate two nn.Linear modules and assign them as
8         member variables.
9     """
10    def __init__(self, input_size, hidden_layers, output_size):
11        super(TwoLayerNet, self).__init__()
12        self.l1 = nn.Linear(input_size, hidden_layers)
13        self.relu = nn.ReLU()
14        self.l2 = nn.Linear(hidden_layers, output_size)
15
16    def forward(self, x):
17        """
18            In the forward function we accept a Tensor of input data and we must return
19            a Tensor of output data. We can use Modules defined in the constructor as
20            well as arbitrary (differentiable) operations on Tensors.
21        """
22        y_pred = self.l1(x)
23        y_pred = self.relu(y_pred)
24        y_pred = self.l2(y_pred)
25
26        return y_pred

```

Hands-on coding & calculating exercises

# Exemplaric Lecture

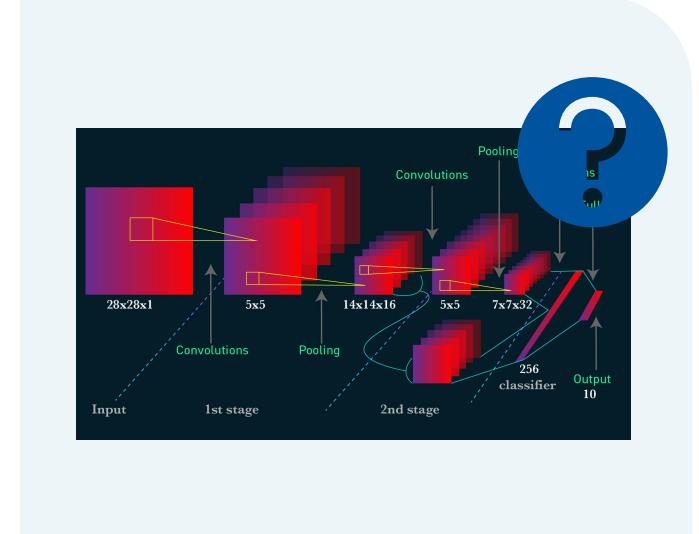
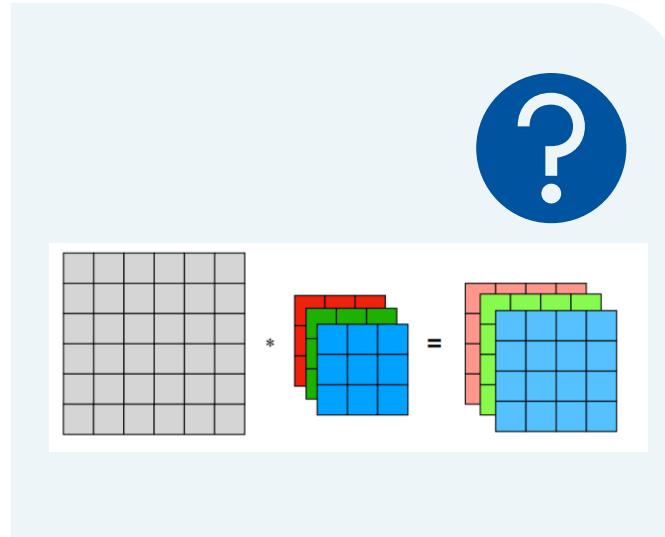
# What we'll learn today



What makes image processing challenging?

How do convolutional filters work?

What do we need to build a Convolutional NN?



# What makes image processing challenging?

**What we see:**



**How its encoded**



Colours are encoded into RGB components.

Input size [pxl]: 32 x 32  
Number of input *channels*: 3  
Notation: 32 x 32 x 3

# What makes image processing challenging?

What we see:



How its encoded

|    |    |    |   |   |   |
|----|----|----|---|---|---|
| 5  | 10 | 2  | 4 | 2 | 0 |
| 3  | 6  | 2  | 0 | 1 | 3 |
| 0  | 7  | 10 | 4 | 8 | 0 |
| 11 | 10 | 8  | 0 | 3 | 7 |
| 10 | 4  | 8  | 6 | 0 | 4 |
| 5  | 12 | 11 | 0 | 2 | 0 |

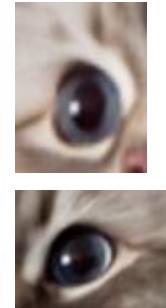
|   |   |    |    |   |   |
|---|---|----|----|---|---|
| 5 | 6 | 2  | 3  | 2 | 0 |
| 5 | 6 | 2  | 0  | 1 | 3 |
| 0 | 5 | 12 | 2  | 5 | 6 |
| 2 | 2 | 8  | 0  | 3 | 7 |
| 0 | 4 | 5  | 10 | 0 | 4 |
| 4 | 1 | 4  | 0  | 2 | 9 |

|    |    |    |   |   |   |
|----|----|----|---|---|---|
| 6  | 10 | 2  | 4 | 2 | 2 |
| 3  | 6  | 12 | 0 | 1 | 6 |
| 0  | 5  | 2  | 4 | 5 | 4 |
| 11 | 13 | 8  | 1 | 3 | 7 |
| 9  | 4  | 4  | 6 | 0 | 9 |
| 5  | 12 | 5  | 8 | 2 | 9 |

Images are matrices.

# What makes image processing challenging?

**What we see:**



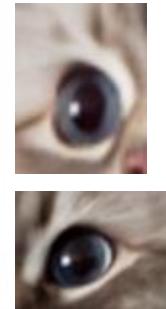
Patterns might repeat among the images.

**Difficulty 1: MLPs do not consider spatial data structure**

- MLPs treat each feature **independently** without considering the spatial structure of the data
- **Spatial arrangement of pixels** is crucial for understanding the image

# What makes image processing challenging?

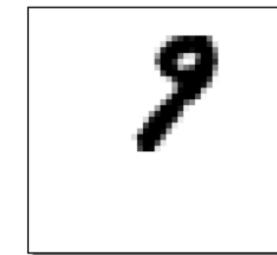
**What we see:**



Patterns might repeat among the images.

## Difficulty 2: Lack of translational invariance

- MLPs are **not robust to shifts** in the image
- If an object is in a **different location**, the MLP may have difficulty recognizing it
- MLP might need to relearn for every feature position



Argue why MLPs are not perfect  
for image recognition

# What makes image processing challenging?



No. of input values is quite high, especially for high-resolution images.

Input size [pxl]:  $32 \times 32$   
Number of input *channels*: 3  
Notation:  $32 \times 32 \times 3$

$32 * 32 * 3 = 3072$  input values

How can we process this using a MLP?

Make an example for parameter

# What makes image processing challenging?



How can we process this using a  
MLP?

10 types of cats &  
use 100 neurons

Model: MLP  
Input: 32 x 32 x 3  
Hidden Layer: 100 nodes  
Output Classes: 10

- Each neuron in the **hidden layer** has
  - 3072 weights (one per input)
  - a bias term
$$3072 * 100 + 100 \text{ parameters}$$
- Each neuron in the **output layer** has
  - 100 weights (one per neuron in the hidden layer)
  - a bias term
$$100 * 10 + 10 \text{ parameters}$$

Total number of parameters:  $(3072 \times 100 + 100) + (100 \times 10 + 10) = 308,310$

Make an example for parameter

# What makes image processing challenging?



How can we process this using a  
MLP?

10 types of cats &  
use 100 neurons

## Difficulty 3: High number of parameters

- Image input vectors can be very large
- Large number of parameters in the network
- Risk of
  - overfitting
  - slower training times



## Caution:

Make sure people know what learnable „*parameters*“ in NNs are i.e. that here you are referring to the „*weights*“.

# What makes image processing challenging?

Difficulty 1: MLPs do not consider spatial data structure

Difficulty 2: Lack of translational invariance

Difficulty 3: High number of parameters

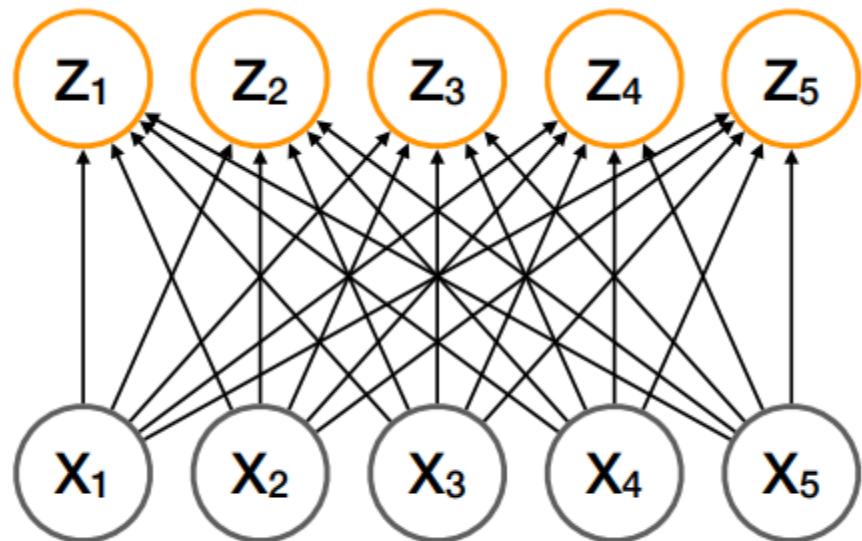


What do we want/need?

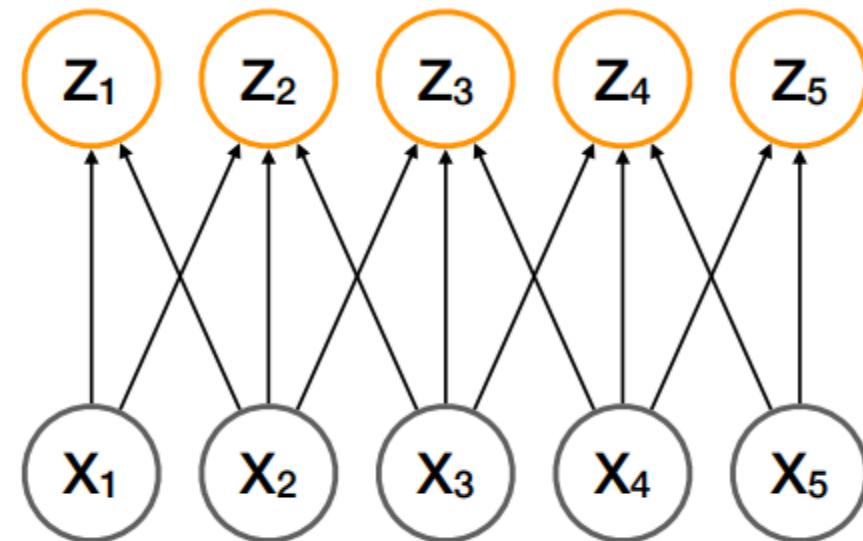
1. **Local connections:** Connections only between neighbouring inputs (pxls)
2. **Translational invariance:**
  - Shifted input results in similar output
  - Parameter sharing

# What makes image processing challenging?

## Solution: Local Connections



**Fully Connected Layer**



**Locally Connected**

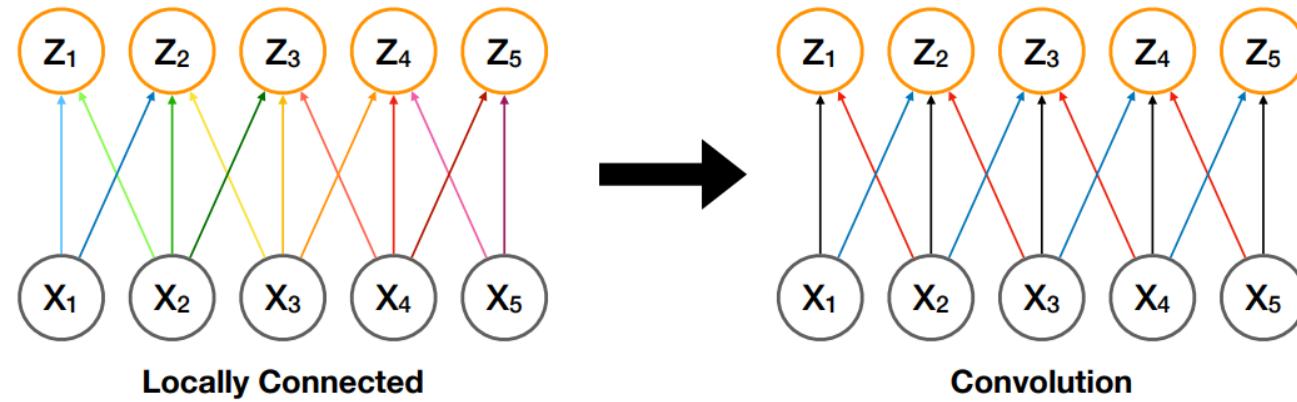
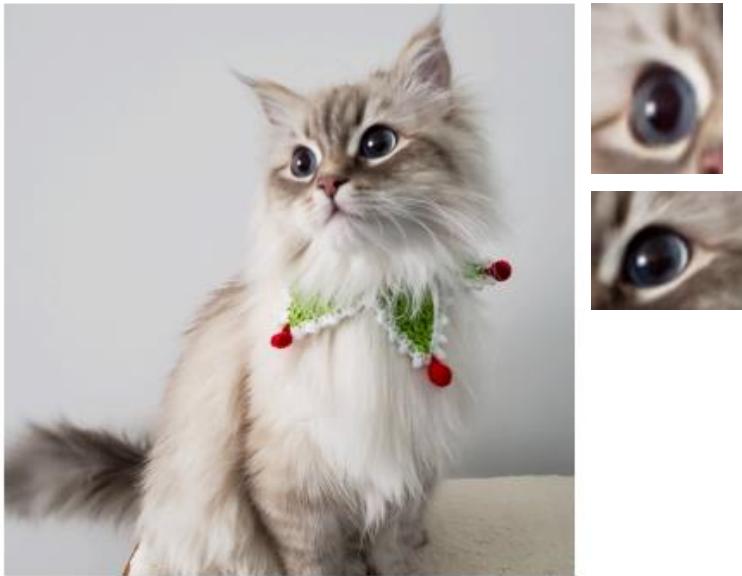


- Image features are (typically) limited to a set of neighboring pixels
- Not entire image needs to be ‘seen’ at once

# What makes image processing challenging?

## Solution: Parameter Sharing

**What we see:**



- Same features can appear at various positions
- Weights that detect features can be repeatedly applied  
-> reduces total no. of parameters

This can be achieved using convolutional filters.

# What makes image processing challenging?

Difficulty 1: MLPs do not consider spatial data structure

Difficulty 2: Lack of translational invariance

Difficulty 3: High number of parameters



CNNs can solve these problems!

CNNs...

- capture spatial structures of images
- are translationally equivariant
- have fewer parameters than MLPs

# How do convolutional filters work?

## Convolution:

Mathematical operation combining two functions to generate a third function.

$$(f * g)(t) \stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau$$

To convolve a kernel with an input signal:  
flip the signal, move to the desired time,  
and accumulate every interaction with the kernel

Image: [cloudzsexy.com](http://cloudzsexy.com)

Depending on the audience you might want to introduce some math.

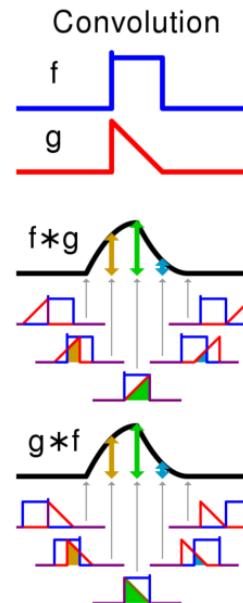
# How do convolutional filters work?

## Convolution:

Mathematical operation combining two functions to generate a third function.

$$(f * g)(t) \stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau$$

To convolve a kernel with an input signal:  
 flip the signal, move to the desired time,  
 and accumulate every interaction with the kernel



- Continuous convolution for real-valued functions  $f$  and  $g$   $(f * g)(t) = \int_{-\infty}^{\infty} f(\tau)g(t - \tau)d\tau$
- Our image inputs are discrete
- Discrete formulation  $(f * g)(n) = \sum_{m=-\infty}^{\infty} f(m)g(n - m)$
- Two dimensional case with Image  $I$  and Kernel  $K$ :  

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n) = (K * I)(i, j)$$

Image: [cloudzsexy.com](http://cloudzsexy.com)

# How do convolutional filters work?

**Image I**

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

$$\begin{array}{c} \text{Image I} \\ \times \\ \text{Kernel K} \\ = \\ \text{I} * \text{K} \end{array}$$

|   |   |   |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n) = (K * I)(i, j)$$

# How do convolutional filters work?

**Image I**

|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 1 | 0 | 1 | 1 | 0 | 0 | 0 |
| 0 | 1 | 2 | 1 | 1 | 0 | 0 |
| 0 | 0 | 0 | 1 | 0 | 1 | 1 |
| 0 | 0 | 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 0 | 0 | 0 | 0 |

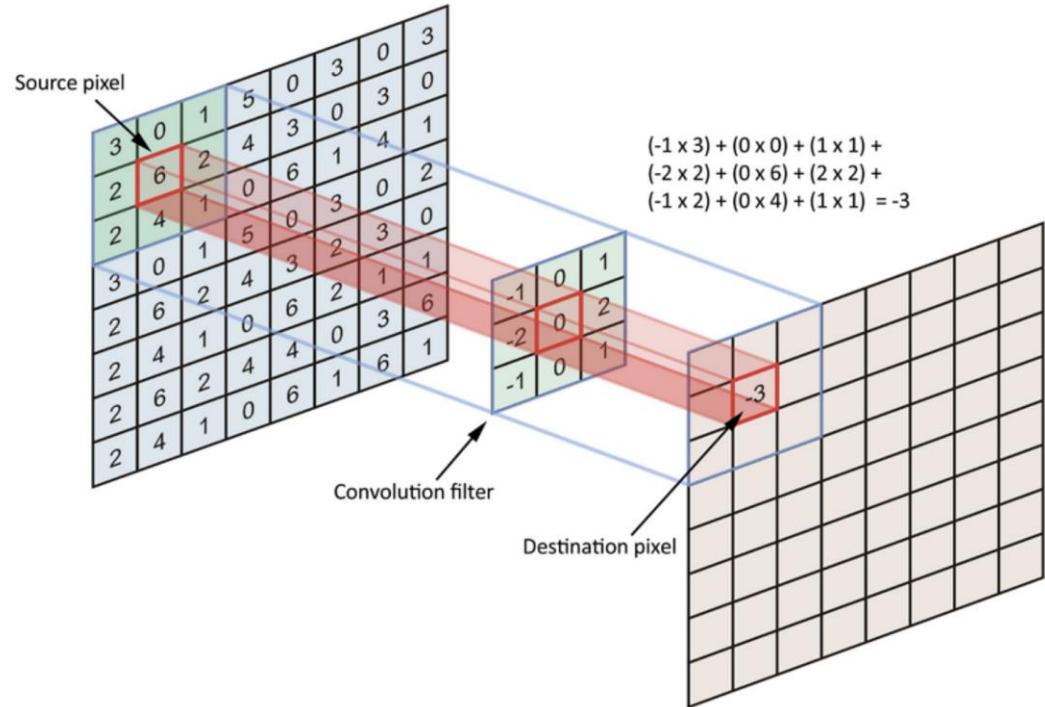
\*      =

|   |   |   |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

|   |  |  |
|---|--|--|
| 4 |  |  |
|   |  |  |
|   |  |  |
|   |  |  |

$$(I * K)(i, j) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} I(m, n)K(i - m, j - n) = (K * I)(i, j)$$

# How do convolutional filters work?



[Image: Simple Introduction to Convolutional Neural Networks](#)  
 | by Matthew Stewart, PhD | Towards Data Science

And show 3D images to make it even clearer.

# How do convolutional filters work?

## Convolutional Filters:

- Used to **extract image features**  
i.e. recognize specific patterns in the input  
-> can be edges, textures or structural elements.
- Filter **weights are learned** during training

Early layers learn **simple features** like edges, later layers combine them into **more complex features** like shapes or objects.

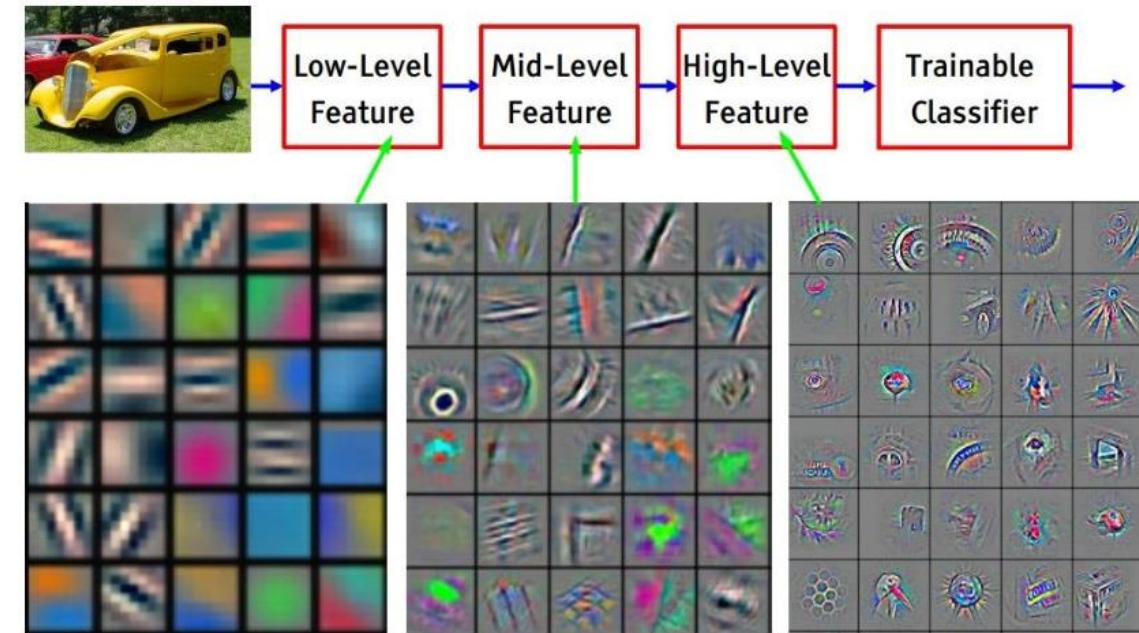
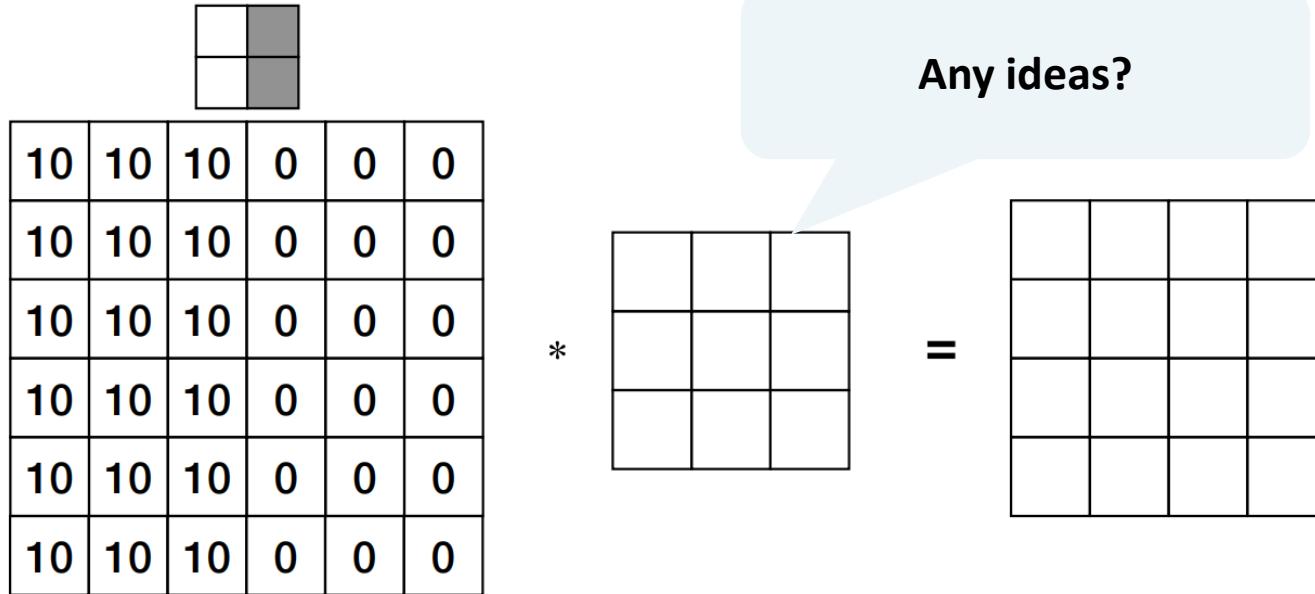


Image: [tpsearchtool.com](http://tpsearchtool.com)

# How do convolutional filters work?

## Application: Edge detection



The diagram illustrates the application of a convolutional filter for edge detection. It shows a 6x6 input matrix and a 3x3 filter kernel. The input matrix has values 10 for most cells and 0 for the last two columns. The filter kernel is a 3x3 matrix with a central value of 1 and other values of 0. A question bubble asks "Any ideas?".

Any ideas?

|    |    |    |   |   |   |
|----|----|----|---|---|---|
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |

\*

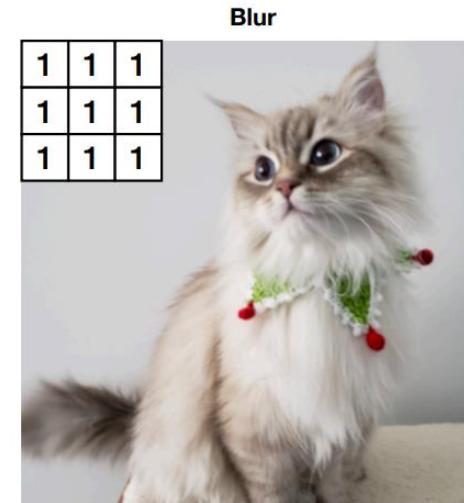
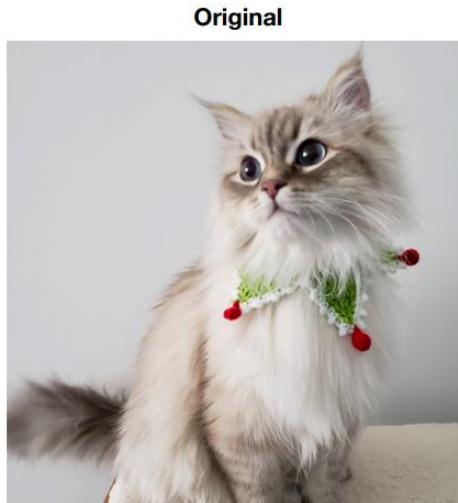
|  |  |  |
|--|--|--|
|  |  |  |
|  |  |  |
|  |  |  |

=

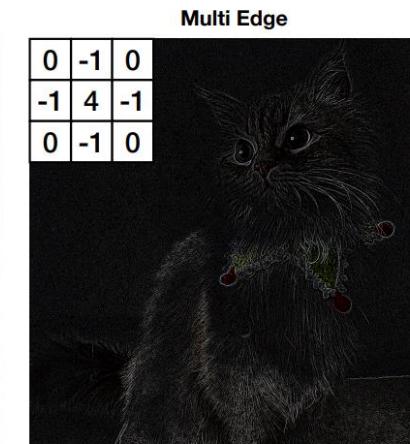
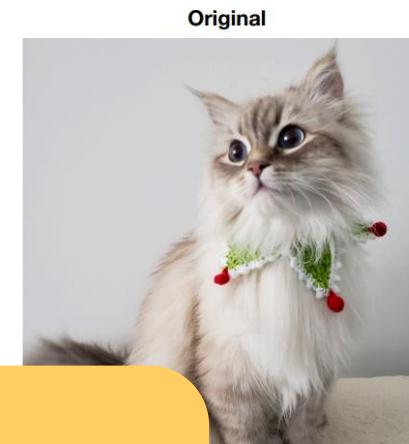
|  |  |  |
|--|--|--|
|  |  |  |
|  |  |  |
|  |  |  |

# How do convolutional filters work?

## Application: Blurring and Sharpening Images



There are many specific filters for specific purposes.



Demonstrate examples of classic filter operations.

# How do convolutional filters work?

## Application: Blurring and Sharpening Images

There are many specific filters for specific purposes.

|    |    |    |
|----|----|----|
| 1  | 1  | 1  |
| 0  | 0  | 0  |
| -1 | -1 | -1 |



How can we generalize?

|       |       |       |
|-------|-------|-------|
| $W_1$ | $W_2$ | $W_3$ |
| $W_4$ | $W_5$ | $W_6$ |
| $W_7$ | $W_8$ | $W_9$ |

$$+ \boxed{B}$$

This is what is learned / updated in our network!

Idea: Replace kernel entries (i.e. hard-coded filter values) with trainable **weights + bias term**  
→ Convolutional Neural Network Layer

Now explain that these filter values are what we want to learn!

# What do we need to build a Convolutional NN?

We can build it ourselves or use a ML library.

Popular:

- Tensorflow
- PyTorch

Building block-like creation of networks.

Important technical details for a Convolutional Layer:

***Channels • Kernel size • Padding • Strides***

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()

Image: Training
```

Make sure students understand common terminology.

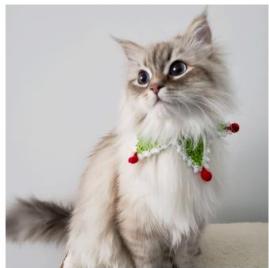
# What do we need to build a Convolutional NN?

Important technical details for a Convolutional Layer:

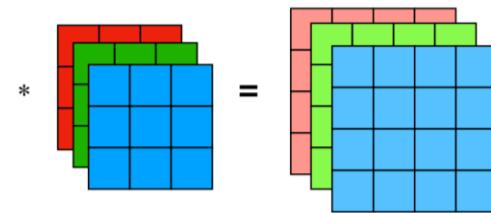
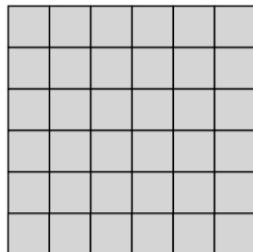
*Channels • Kernel size • Padding • Strides*

## Channels

What we see:



How its encoded



One filter may not be sufficient to extract all important features! Example: RGB images

- Use multiple filters for multiple output channels/feature maps
- Each filter has independent weights

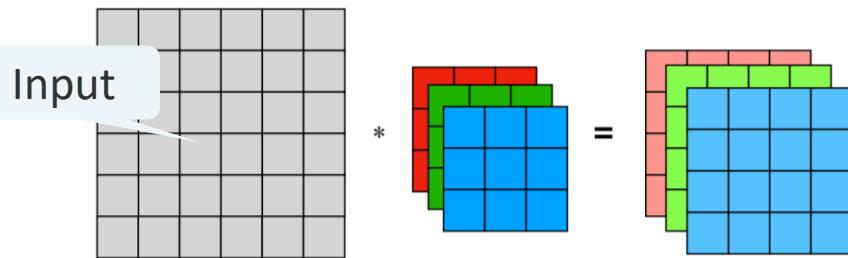
Make sure students understand common terminology.

# What do we need to build a Convolutional NN?

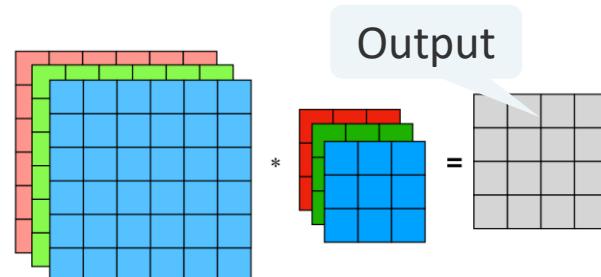
Important technical details for a Convolutional Layer:

*Channels • Kernel size • Padding • Strides*

## Channels



One filter per input channel, sum over each convolution result



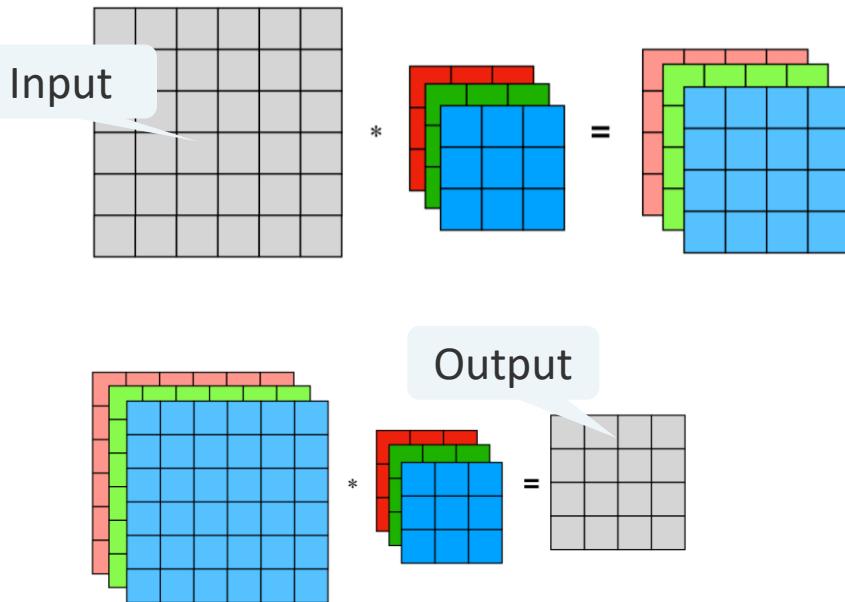
We typically only want a single output channel, containing information from all inputs

# What do we need to build a Convolutional NN?

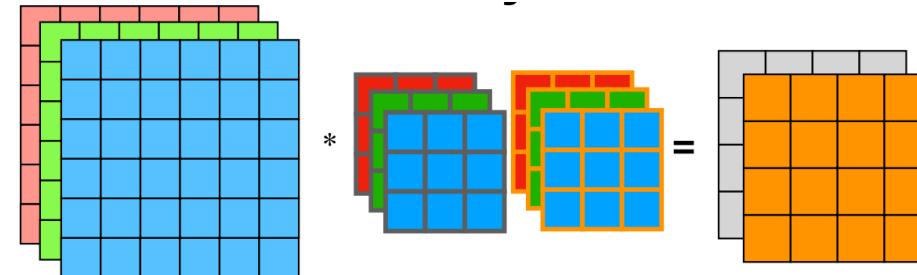
Important technical details for a Convolutional Layer:

*Channels • Kernel size • Padding • Strides*

## Channels



However, multiple input and output channels are also possible!

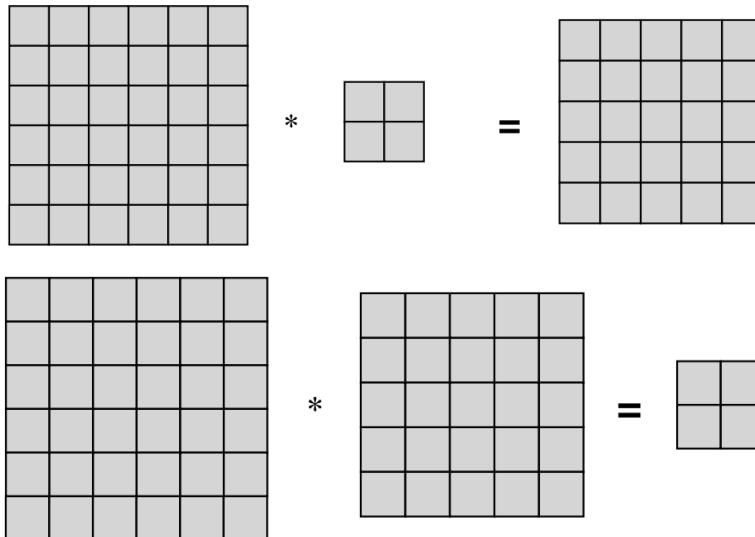


# What do we need to build a Convolutional NN?

Important technical details for a Convolutional Layer:

*Channels • Kernel size • Padding • Strides*

## Kernel Size / Filter Size:



Kernel size dictates “how far” pixels are connected (in their neighborhood):

- Too small kernel: insufficient to cover interesting features
- Too large kernel: results in too many weights

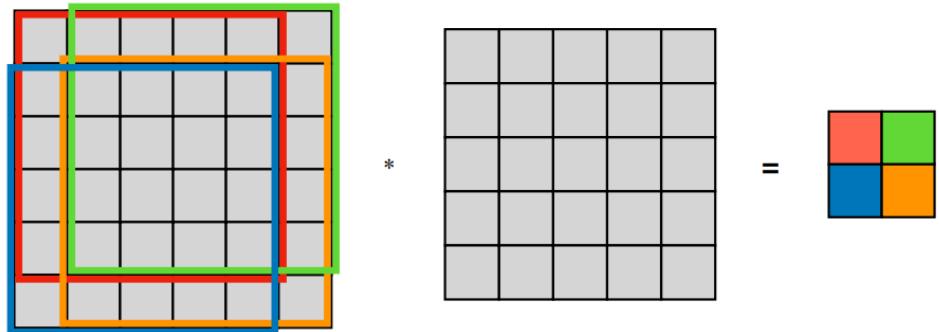
Make sure students understand common terminology.

# What do we need to build a Convolutional NN?

Important technical details for a Convolutional Layer:

*Channels • Kernel size • Padding • Strides*

## The Shrinking of Convolution Outputs



Kernel needs to overlap fully with input image

For kernel size > 1:  
reduces the output size (compared to input)  
→ default setting (“*valid padding*”)

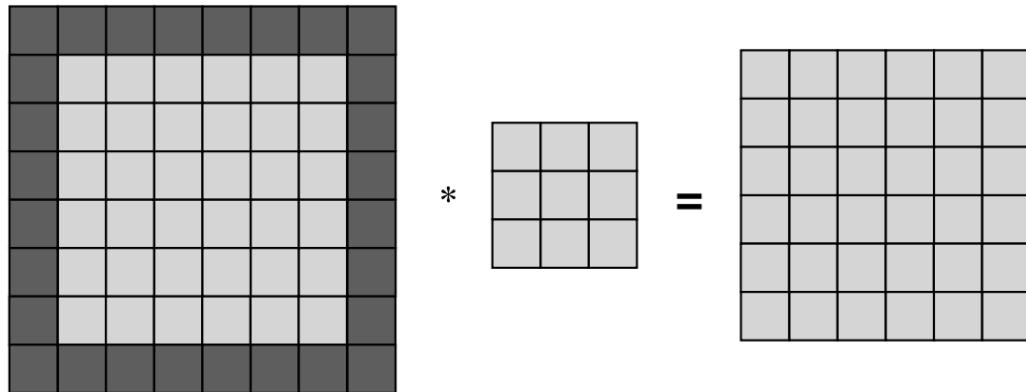
For understanding padding it  
helps to see an image „shrink“.

# What do we need to build a Convolutional NN?

Important technical details for a Convolutional Layer:

*Channels • Kernel size • Padding • Strides*

## Padding



‘Same’ padding: pads such that input and output are same size

Kernel “cannot” be applied to pixels at the edges so that our output images preserves its image size

-> Padding is a solution for artificially enlarging our images, so that our images do not “shrink” during convolution

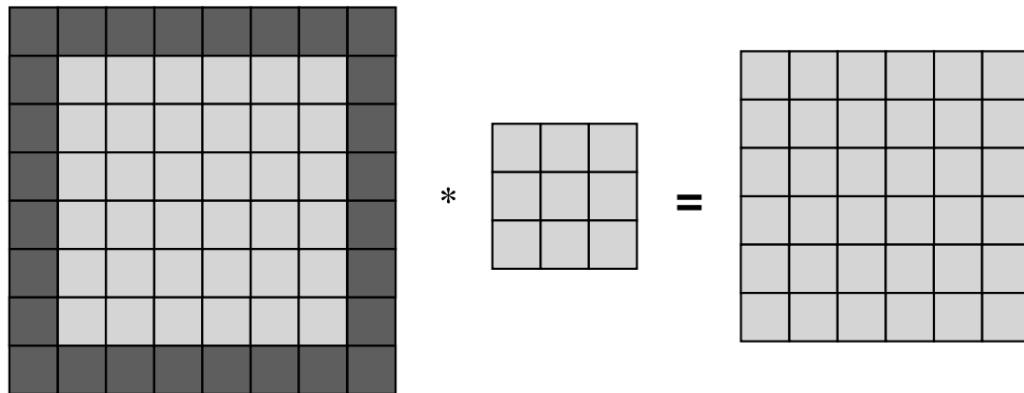
# What do we need to build a Convolutional NN?

Important technical details for a Convolutional Layer:

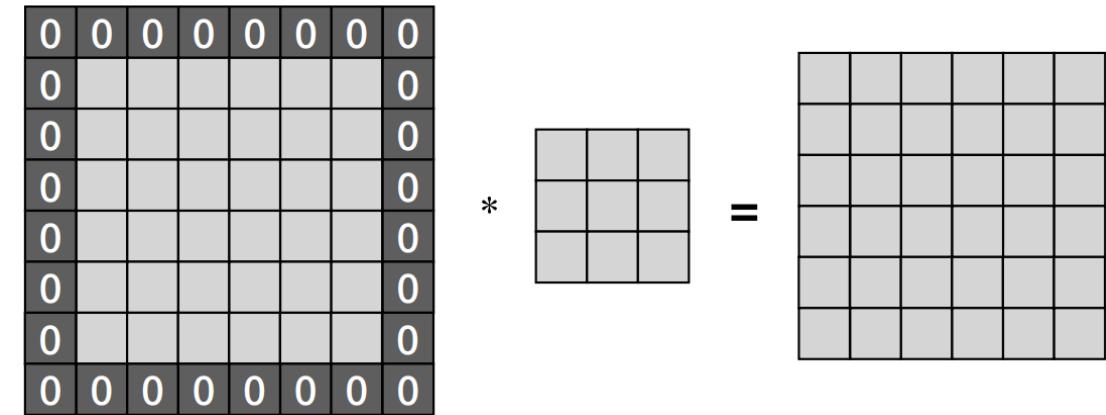
*Channels • Kernel size • Padding • Strides*

There are many more padding methods. [Guide to Different Padding Methods for CNN Models - \(analyticsindiamag.com\)"](#)

## Padding



'Same' padding: pads such that input and output are same size



'Zero' padding: pads given area using 0 as a fill value  
Can have 'Zero' (value) and 'Same' (shape) padding

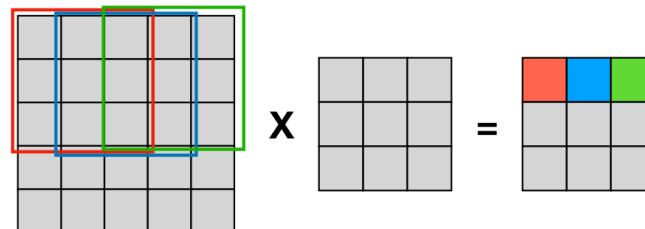
# What do we need to build a Convolutional NN?

Important technical details for a Convolutional Layer:

*Channels • Kernel size • Padding • Strides*

## Stride

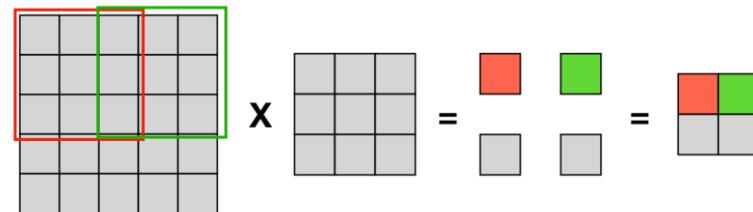
Stride = 1: Move kernel by 1 pixel



Strides describes how the convolutional kernel moves (or “jumps”) in each step

Strides > 1 act as *downsampling*.

Stride = 2: Move kernel by 2 pixel



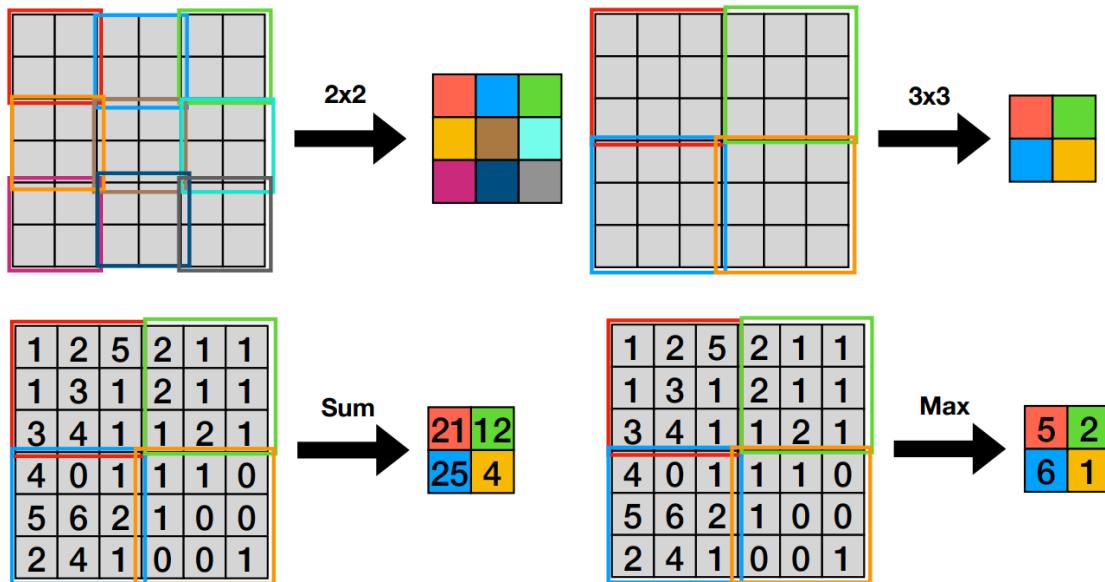
Let them calculate examples on a sheet of paper.

# What do we need to build a Convolutional NN?

Important technical details for a Convolutional Layer:

*Channels • Kernel size • Padding • Strides*

## Pooling



Reducing the image dimensionality is often required. Options:

- Stride convolutions
- Pooling

### Pooling:

Combines information of several pixels into one pixel

-> for ex. by taking the **sum, mean or max**

Again: Pooling kernel size influences the output size, i.e. how much the image is downsampled

# What do we need to build a Convolutional NN?

```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```

- Readily implemented in most ML libraries
- Important technical details summary:
  - **Channels:** number of input and output channels
  - **Kernel size:** choose appropriate to feature size
  - **Padding:** adds ‘empty’ pixels to manipulate output size
  - **Strides:** ‘step size’, allows for downsampling

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

```
net = Net()
```

[Image: Training](#)

Always come back to your  
overview slide ...

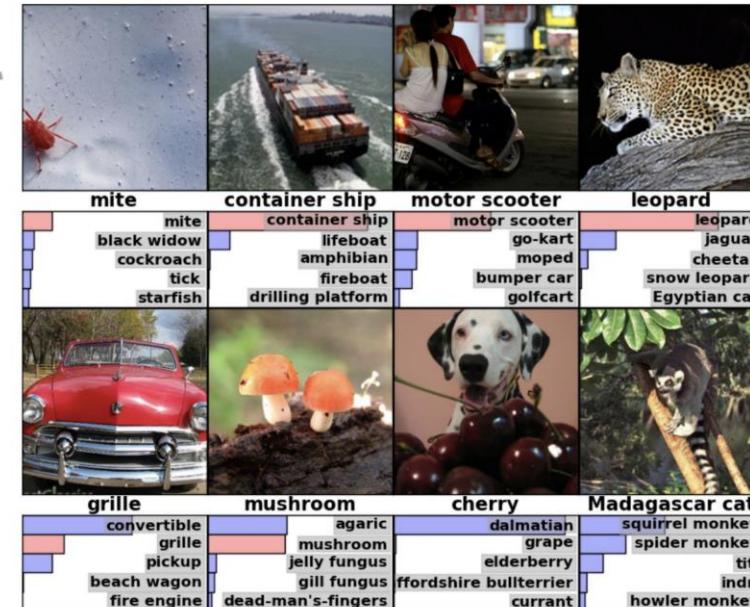
# (Historically) Popular Network Architectures

Many of them were created for the...

## ImageNet Challenge

IMAGENET

- 1,000 object classes (categories).
- Images:
  - 1.2 M train
  - 100k test.



[Image: NLP's ImageNet moment has arrived \(thegradient.pub\)](https://nlp.stanford.edu/pubs/imagenet_moment.pdf)

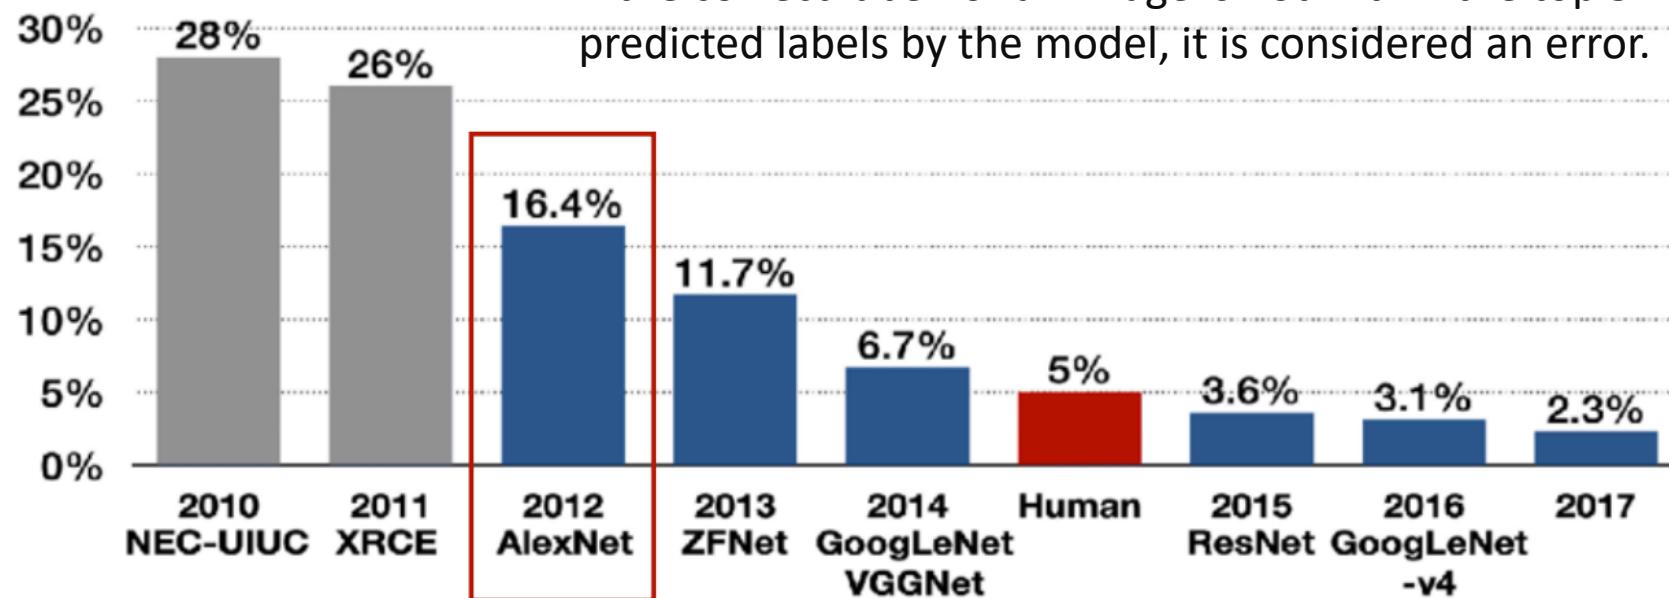


[Image: Meme Creator](#)

# (Historically) Popular Network Architectures

## Image Net Challenge

Top-5 error

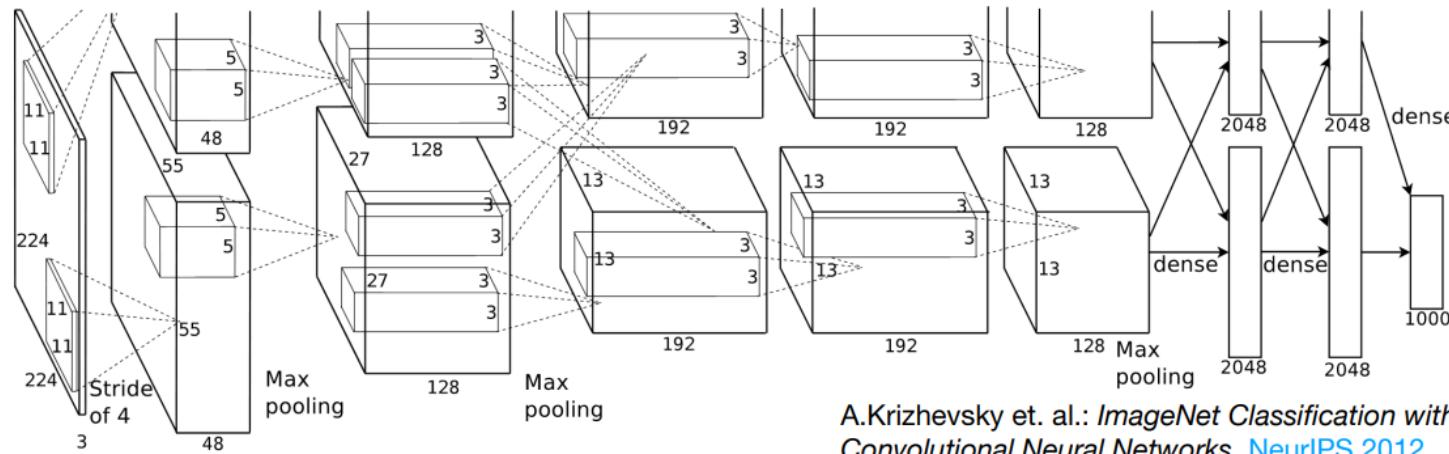


Dae-Young Kang et. al.: Application of Deep Learning  
in Dentistry and Implantology, [10.32542](https://doi.org/10.32542)

# (Historically) Popular Network Architectures

## AlexNet

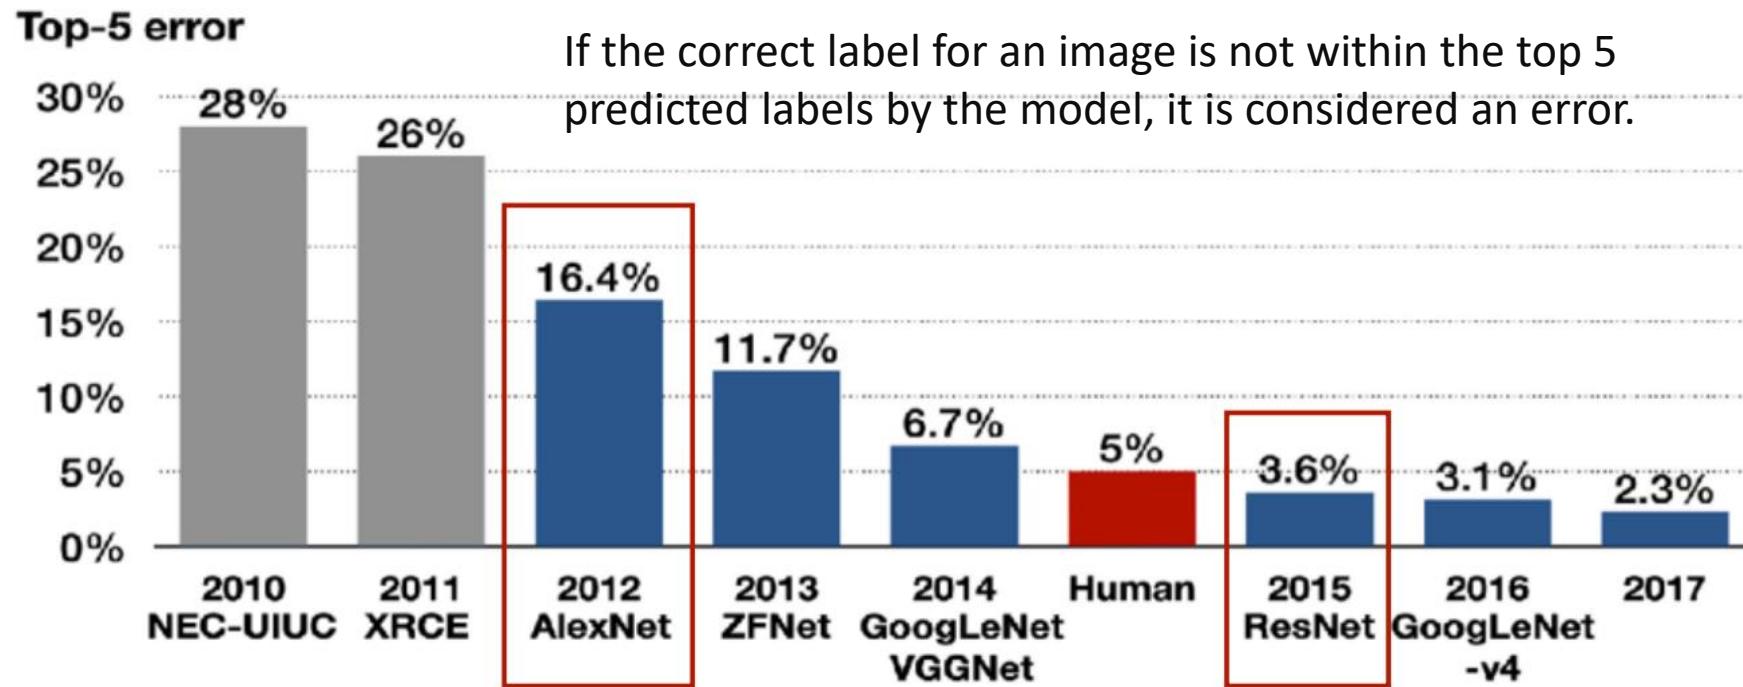
- Archetypal for most convolutional networks
  - Convolutional segment
  - Dense segment



A.Krizhevsky et. al.: *ImageNet Classification with Deep Convolutional Neural Networks*, NeurIPS 2012

# (Historically) Popular Network Architectures

## Image Net Challenge



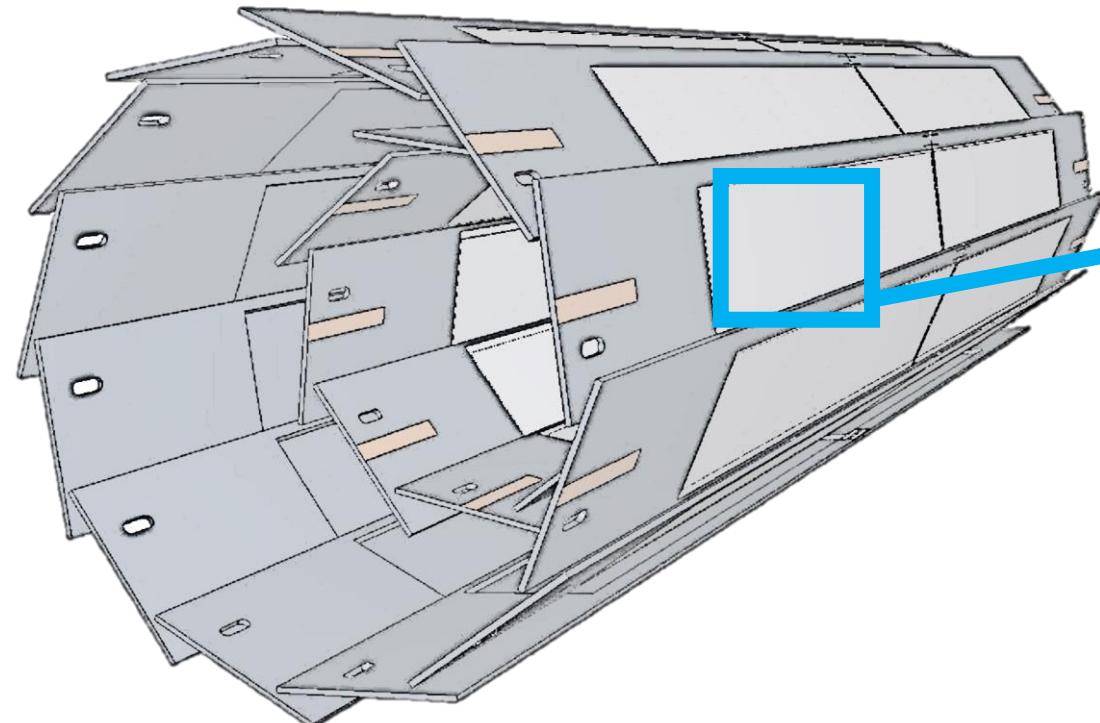
Dae-Young Kang et. al.: Application of Deep Learning  
in Dentistry and Implantology, [10.32542](#)



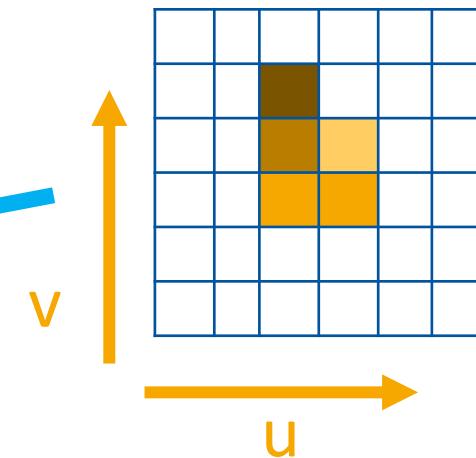
Make sure they understand that **deeper is not always better!**

# (Historically) Popular Network Architectures

...and why they are sometimes not well-suited.



belle2.desy.de



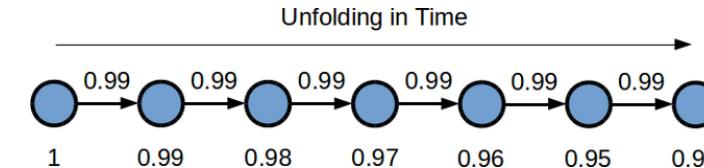
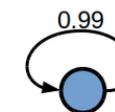
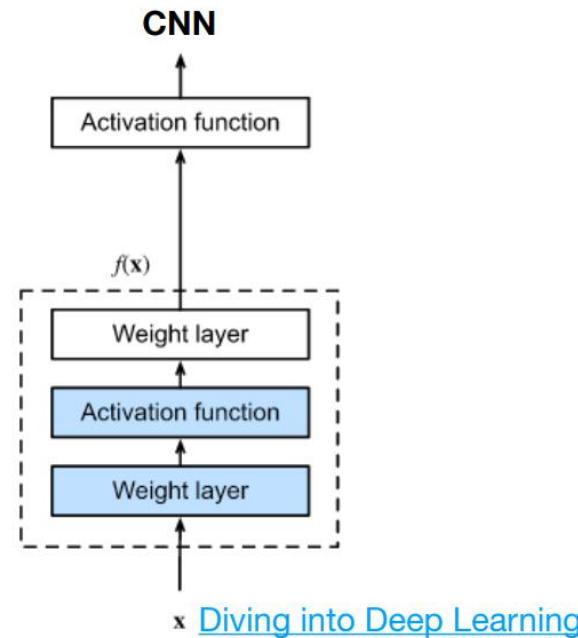
pixel color =  
deposited charge

If we apply a **deep architecture**,  
small images might **disappear**!

# (Historically) Popular Network Architectures

## ResNet & Vanishing Gradient

- Standard CNNs have limited depth, too deep
- Vanishing gradients
  - Negative weights
  - Negative output
  - ReLU maps to zero
  - No gradient
- More likely to occur for deep networks



### Vanishing Gradient Problem

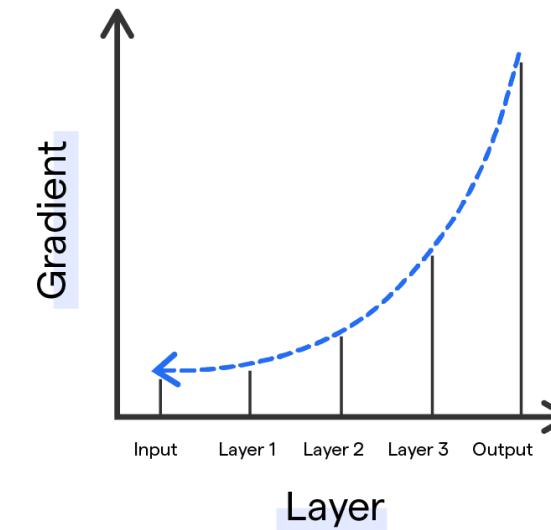


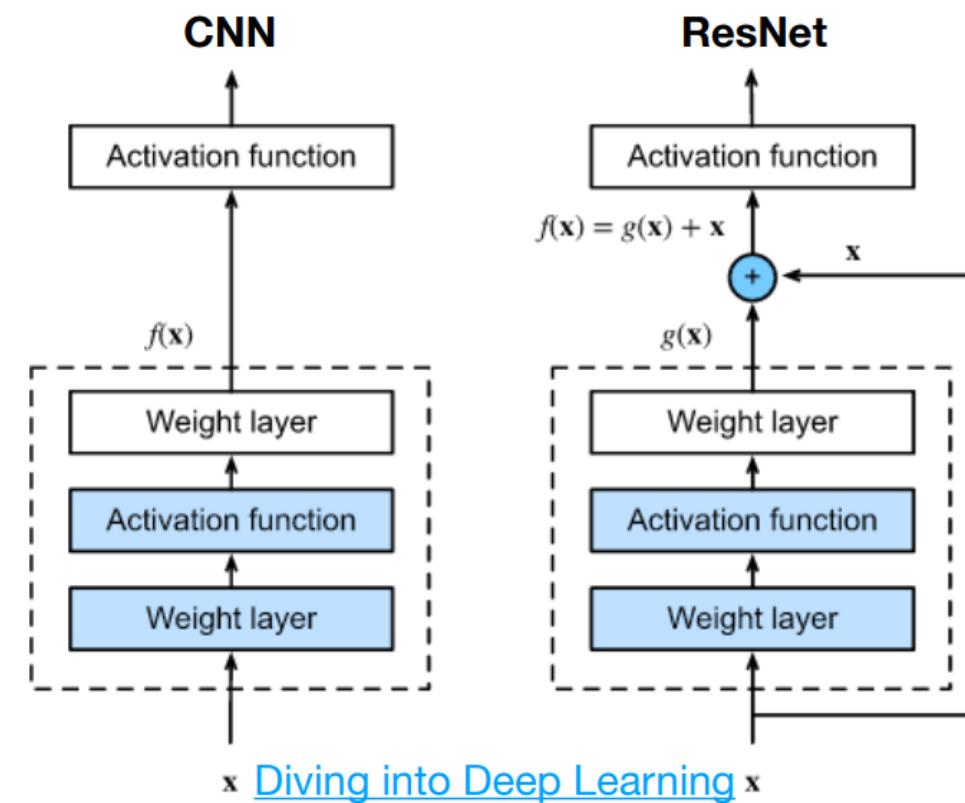
Image: BotPenguin

Do not forget the vanishing gradient problem!

# (Historically) Popular Network Architectures

## ResNet

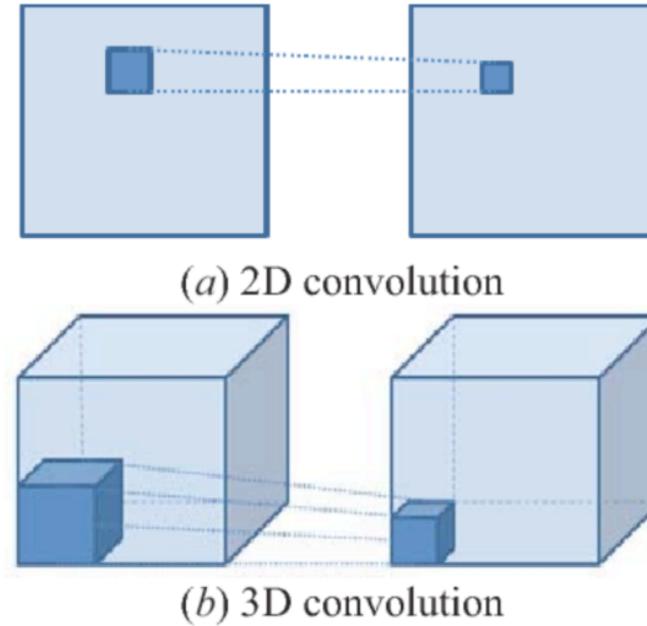
- Standard CNNs have limited depth, too deep  
 → Vanishing gradients
- Solution:
  - Residual/skip connections
  - Gradients can propagate through network easily
  - Conv. layers only learn modification to input



# Outlook: 3D Convolutions

Convolutions not limited to 2D

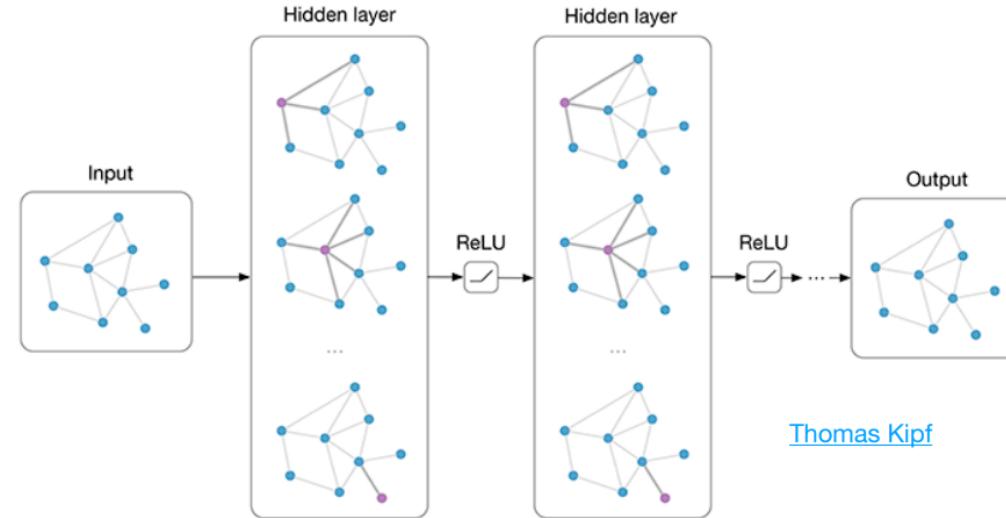
- Concept can be extended to 3D data as well
  - 3D output, 3D kernel, 3D strides, ....
  - Volumetric data processing
- Similar for 1D
  - Time sequence processing



Fan et. al., Lung nodule detection based on  
3D convolutional neural networks, [document/88](#)

# Outlook: Graph Convolutions

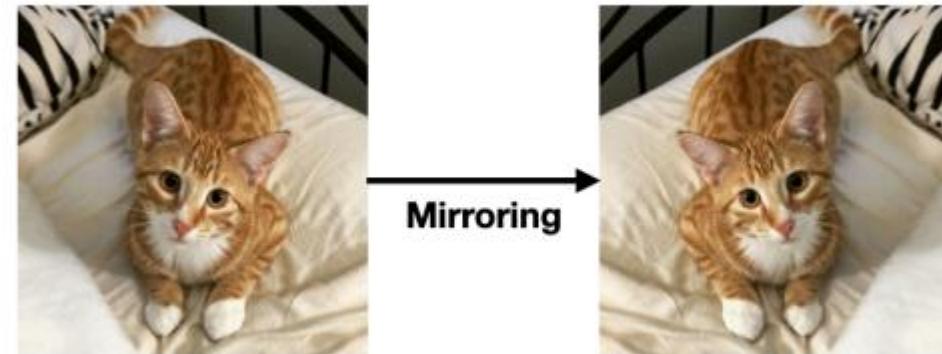
- Graph CNNs
- Operate on graphs with arbitrary node positions
- Do not require fixed structures like CNNs
- More in advanced course



[Thomas Kipf](#)

# Explaining the Training Process

# What do we need to train a Convolutional NN?



## Data Augmentation:

Mirroring, Flipping, Rotating, Cropping of our images etc.

- > artificial increase of available data
- > less overfitting

# What do we need to train a Convolutional NN?

Hyperparameters?

| Hyperparameter         | Range                   | Optimal Value |
|------------------------|-------------------------|---------------|
| learning rate          | [0.001, 0.01, 0.1, 0.2] | 0.01          |
| batch size             | [5, 10, 20]             | 20            |
| epochs                 | [8, 16, 32]             | 32            |
| optimization algorithm | [SGD, RMSprop, Adam]    | SGD           |

[The range and optimal values of hyperparameters for CNN.](#) | Download Scientific Diagram ([researchgate.net](#))

- Fitting a model is an **optimization problem**
- We **cannot test all possibilities**
- We rely on **heuristics** for choosing them

# What do we need to train a Convolutional NN?

Explain tradeoff between training time and quality of the results.

| Hyperparameter         | Range                   | Optimal Value |
|------------------------|-------------------------|---------------|
| learning rate          | [0.001, 0.01, 0.1, 0.2] | 0.01          |
| batch size             | [5, 10, 20]             | 20            |
| epochs                 | [8, 16, 32]             | 32            |
| optimization algorithm | [SGD, RMSprop, Adam]    | Adam          |



Adam Optimizer is often a good choice.

[The range and optimal values of hyperparameters for CNN. | Download Scientific Diagram \(researchgate.net\)](#)



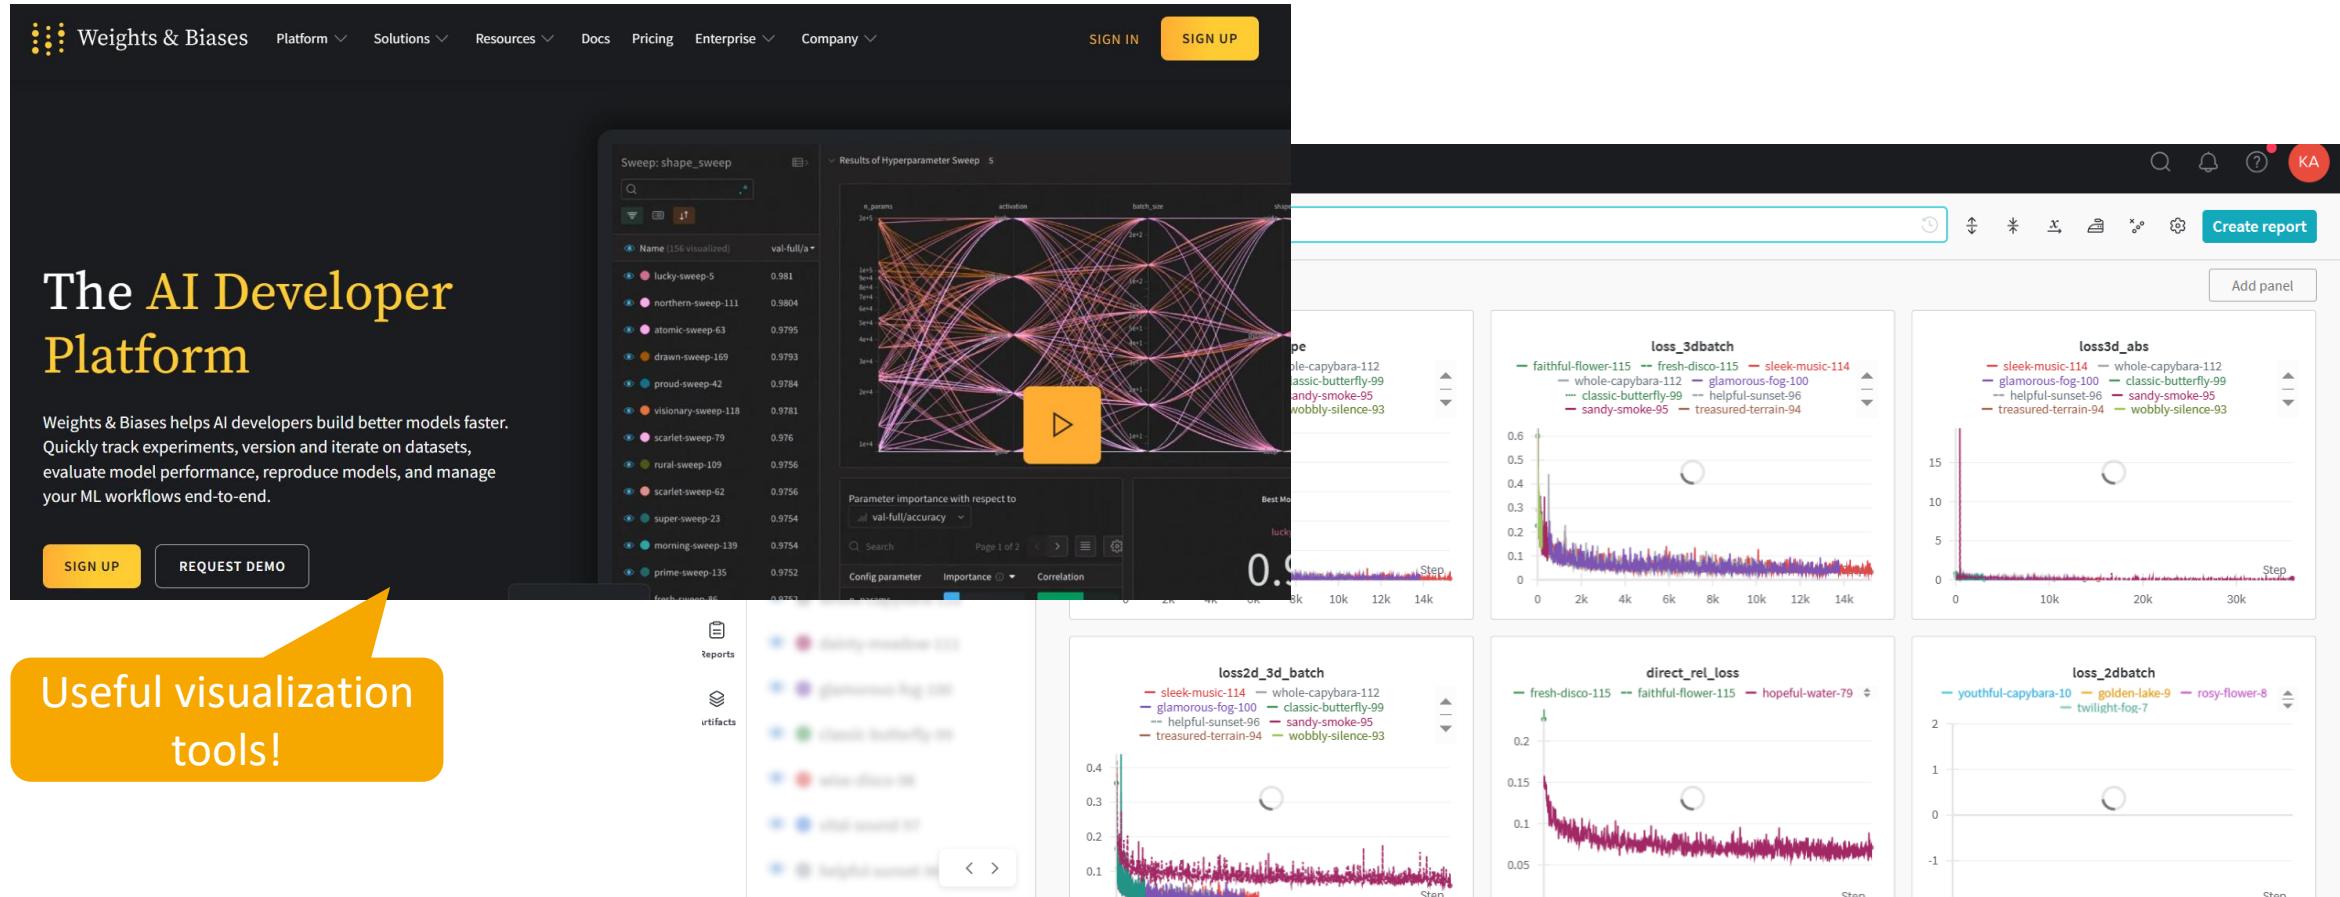
# What do we need to train a Convolutional NN?

The AI Developer Platform

Weights & Biases helps AI developers build better models faster. Quickly track experiments, version and iterate on datasets, evaluate model performance, reproduce models, and manage your ML workflows end-to-end.

[SIGN UP](#) [REQUEST DEMO](#)

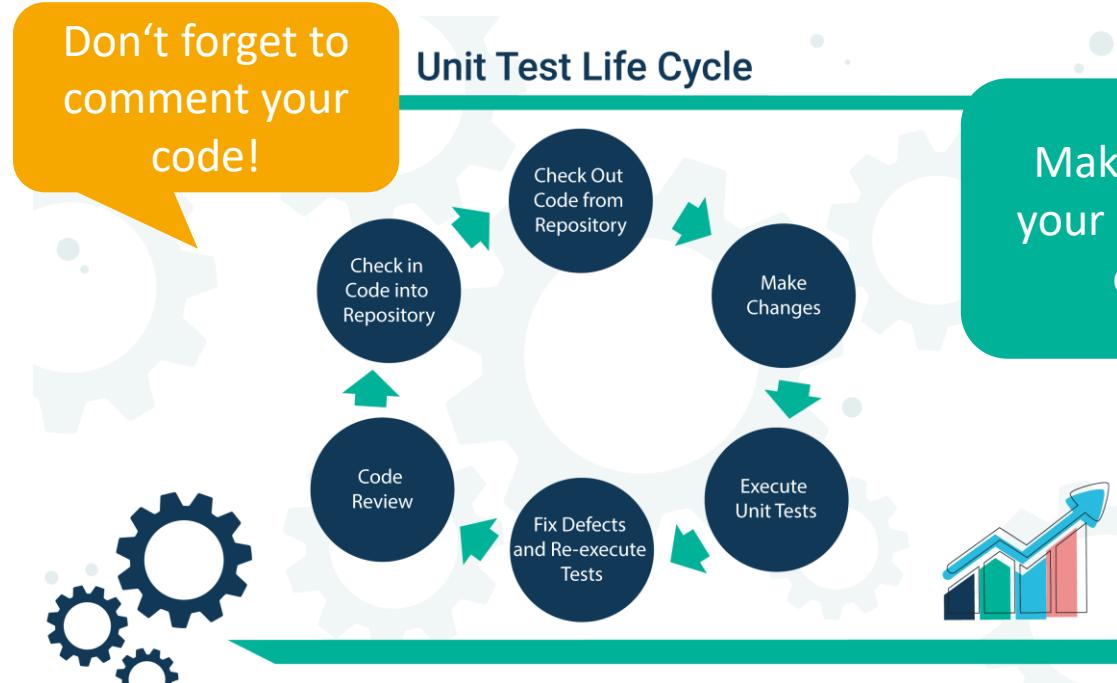
Useful visualization tools!



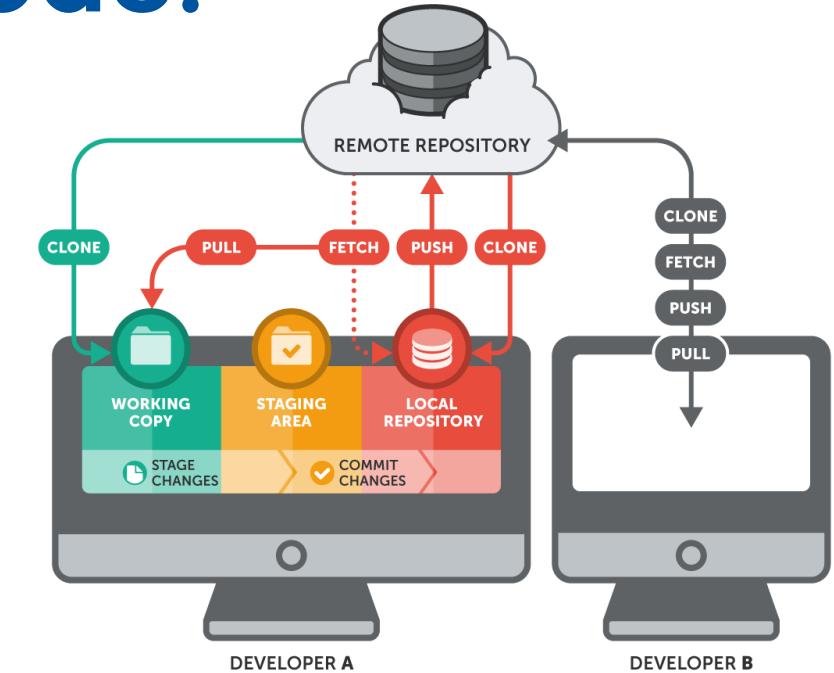
The screenshot displays the Weights & Biases platform interface. At the top, there's a navigation bar with links for Platform, Solutions, Resources, Docs, Pricing, Enterprise, and Company, along with SIGN IN and SIGN UP buttons. Below the navigation is a search bar and a 'Create report' button. The main area features several visualization panels:

- Sweep: shape\_sweep**: A hyperparameter sweep visualization showing connections between parameters like n\_params, activation, batch\_size, and shape.
- Results of Hyperparameter Sweep**: A network graph of parameter interactions.
- Parameter importance with respect to val-full/accuracy**: A line chart showing the importance of various parameters over 14k steps.
- Best Model**: A chart showing accuracy over 14k steps, with a highlighted peak at step 0.
- loss\_3dbatch**, **loss3d\_abs**, **loss2d\_3d\_batch**, **direct\_rel\_loss**, and **loss\_2dbatch**: Multiple line charts showing loss metrics over training steps, each with a legend for different experiment runs.
- reports** and **artifacts**: Side panels for managing project reports and artifacts.

# What do we need to train a Convolutional NN? Clean code.



Make sure to keep your code clean and organized.



[10 Best Git Tutorials for Beginners \[2024 FEB\]-Learn Git online | Quick Code \(medium.com\)](#)

[Image: Unit Test: What is, Advantages & How works? - CodenBox AutomationLab](#)

Use unit tests and git for version control.

# Possible Exercises

# CNN-Theory Tasks



Ask for

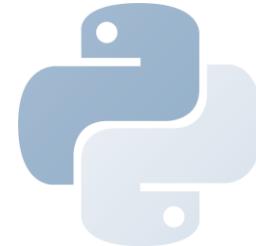
- The number of **learnable parameters**
- The **output size** of a convolutional layer
- The influence of **stride/pooling/padding** etc.
- Setting up a **feasible layer structure**  
(i.e. not too many layers if image is small, add dropouts if no of parameters is high etc.)
- Explanation of **code fragments**
- The influence of **skip connections**
- Benefits of **ReLU**
- Role of **FCN layers** in CNN architectures
- ...

# Coding Tasks



**Refresh Skills**

-> provide tutorial



**Use Python**  
(Julia and R are  
harder to debug)



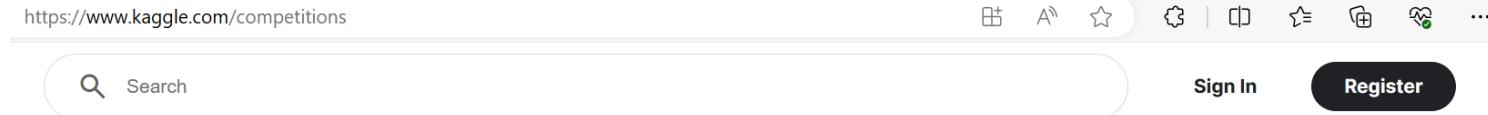
**Either PyTorch  
or Tensorflow**



**Support  
best practices**  
(Git, Wandb / TensorBoard,  
Unit Tests)

# Coding Tasks

- Start with simple **architectures** like LeNet
- **Gradually increase complexity** by adding argumentation, regularization etc.
- Make use of **well-known baselines** (ImageNet etc.)
  
- **Cooperative** tasks
- Set up a **challenge or project**



The screenshot shows a web browser displaying the Kaggle website at <https://www.kaggle.com/competitions>. The page features a search bar and navigation links for 'Sign In' and 'Register'. The main content area is titled 'Competitions' and includes a sub-section for 'Community Competitions'.

## Competitions

Grow your data science skills by competing in our exciting competitions. Find help in the [documentation](#) or learn about [Community Competitions](#).

[Host a Competition](#)



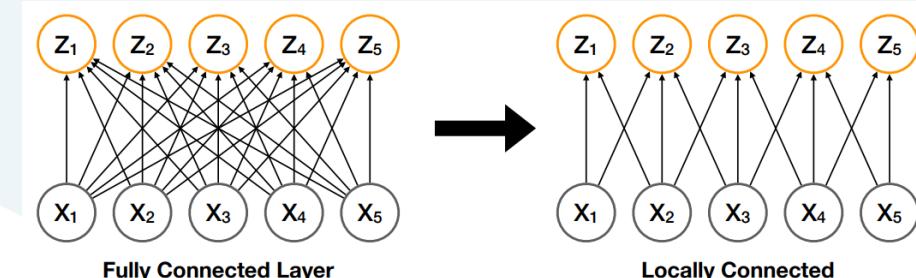
# What might be hard to understand?



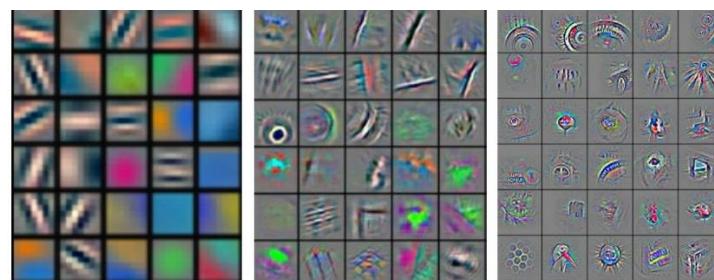
Images as input



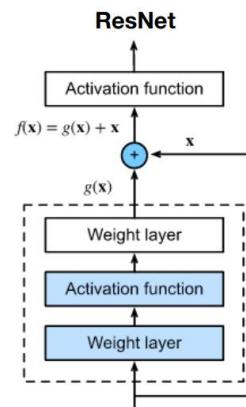
Concept of channels



Reasons for parameter reduction



What filters actually learn



Issues during the training process

Regularization & training tricks

How to set hyperparameters?